

# UNCERTAINTY QUANTIFICATION IN SCIENTIFIC MACHINE LEARNING

*Uncertainty in Scientific Machine Learning from PDEs to Molecules*

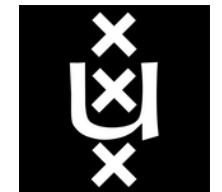
---

Advanced School on Foundation Models for  
Scientific Discovery

*The Abdus Salam International Centre for  
Theoretical Physics, Italy*

DARIO COSCIA

SISSA, University of Amsterdam



# LECTURE OUTLINE 2/2

- **Metrics for Uncertainty Quantification**

- Scoring Rules
- Calibration and Sharpness
- Correlation

- **Uncertainty Quantification for Sequence Models**

- Introduction to Autoregressive Neural Operators
- Functional Generative Networks, PDE-Refiner, BARNN

- **Uncertainty Quantification for Message Passing Network**

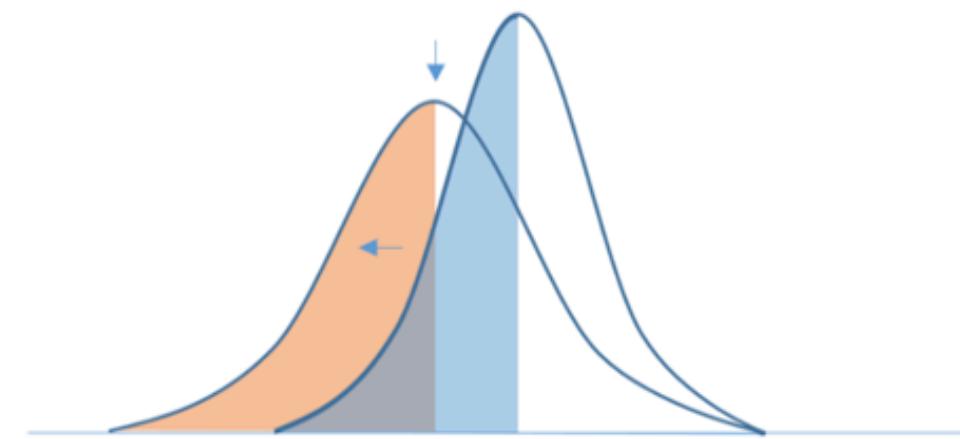
- Introduction to Graph Networks and Machine Learning Interatomic Potentials
- UQ in Machine Learning Interatomic Potentials

# SHORT RECAP...

**Take Aways:** 1<sup>st</sup> There are two types of uncertainty: irreducible uncertainty (aleatoric) and model uncertainty (epistemic).

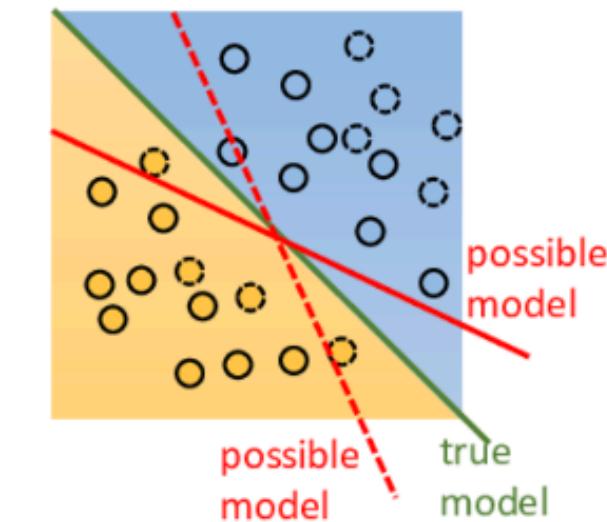
**Aleatoric Uncertainty**

$$y_{\text{data}} = \mathcal{M}_{\text{true}}(x) + z \quad , z \sim p(z)$$



**Epistemic Uncertainty**

$$\mathcal{M}_{\text{model}}(x | \omega) \quad , \omega \sim p(\omega)$$



# SHORT RECAP...

**Take Aways:** 2<sup>nd</sup> The Bayesian approach enables modelling both epistemic and aleatoric uncertainties

$$p(y \mid x) = \int p(y \mid x, \omega)p(\omega \mid \mathcal{D})d\omega \quad \text{with,} \quad p(\omega \mid \mathcal{D}) = \frac{p(\mathcal{D} \mid \omega)p(\omega)}{p(\mathcal{D})}$$

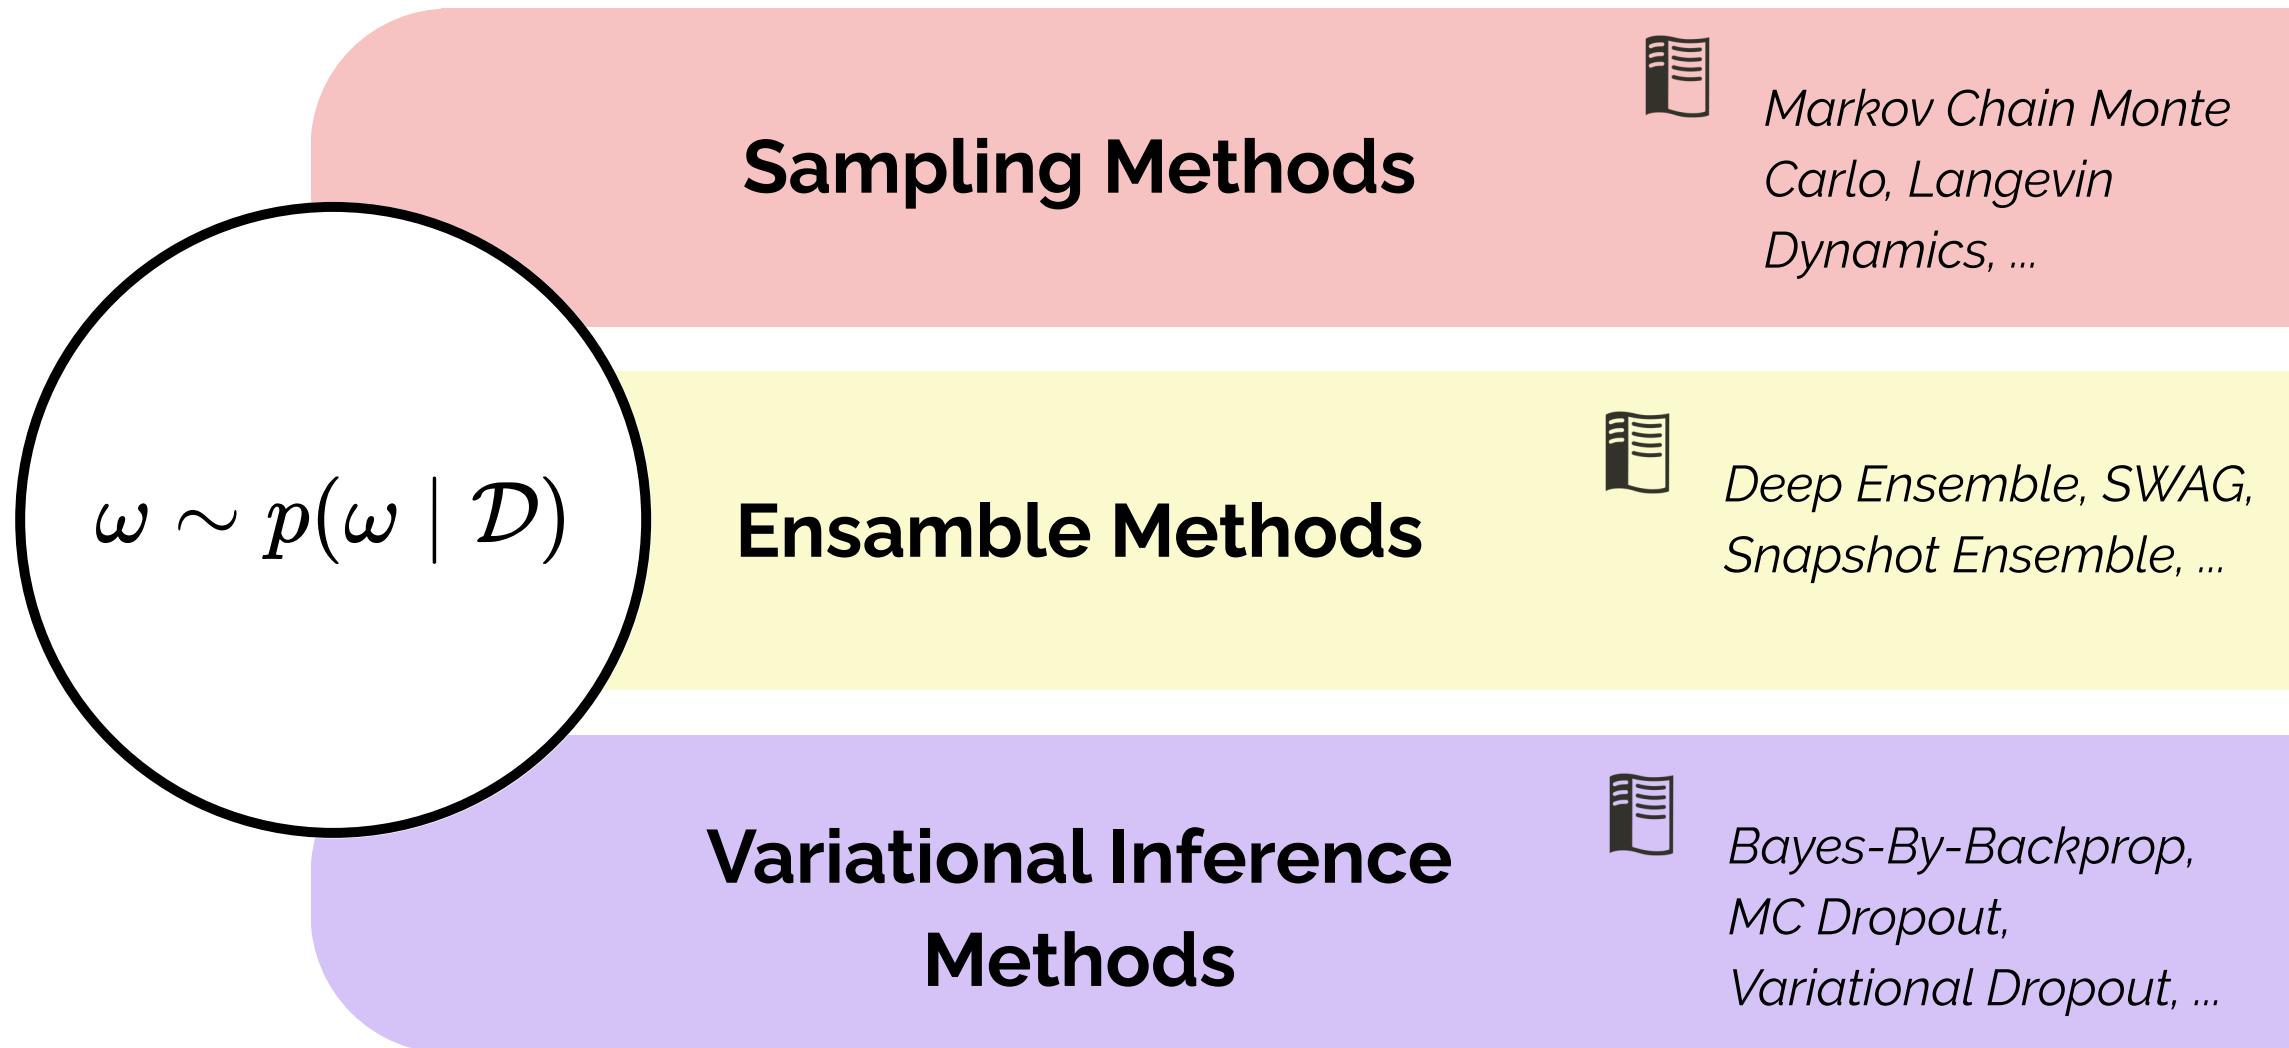
Using the **law of total variance** we can decompose the uncertainty

$$\text{Var}[y \mid x] = \mathbb{E}_{\omega \sim p(\omega \mid \mathcal{D})} [\text{Var}[y \mid x, \omega]] + \text{Var}_{\omega \sim p(\omega \mid \mathcal{D})} [\mathbb{E}[y \mid x, \omega]]$$

*aleatoric*                                   *epistemic*

# SHORT RECAP...

**Take Aways:** 3<sup>rd</sup> Many ways to compute the posterior distribution



# A SMALL NOTE ON LIKELIHOOD FUNCTION

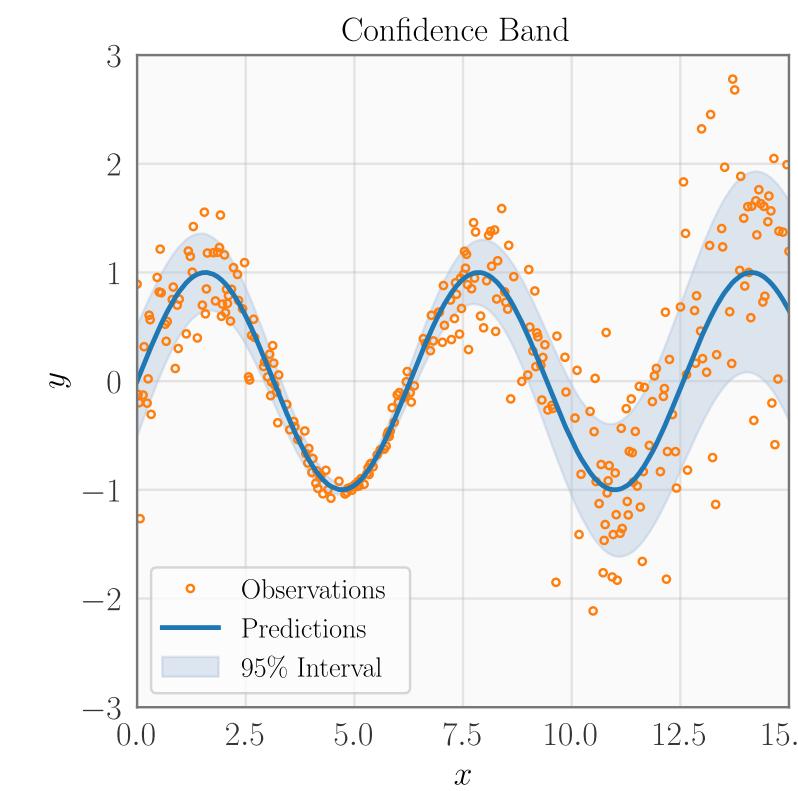
**Setting:** We are given a dataset  $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$  of  $N$  observations, and we want a model such that  $y_i = \phi_\omega(x_i)$  (regression task), **which likelihood do I choose?**

**Gaussian model:**  $p(\mathcal{D} | \omega) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(y_i - \phi_\omega(x_i))^2}{2\sigma^2}} \implies -\nabla_\omega \log p(\mathcal{D} | \omega) \propto \nabla_\omega \text{MSE}(y, \phi_\omega(x))$

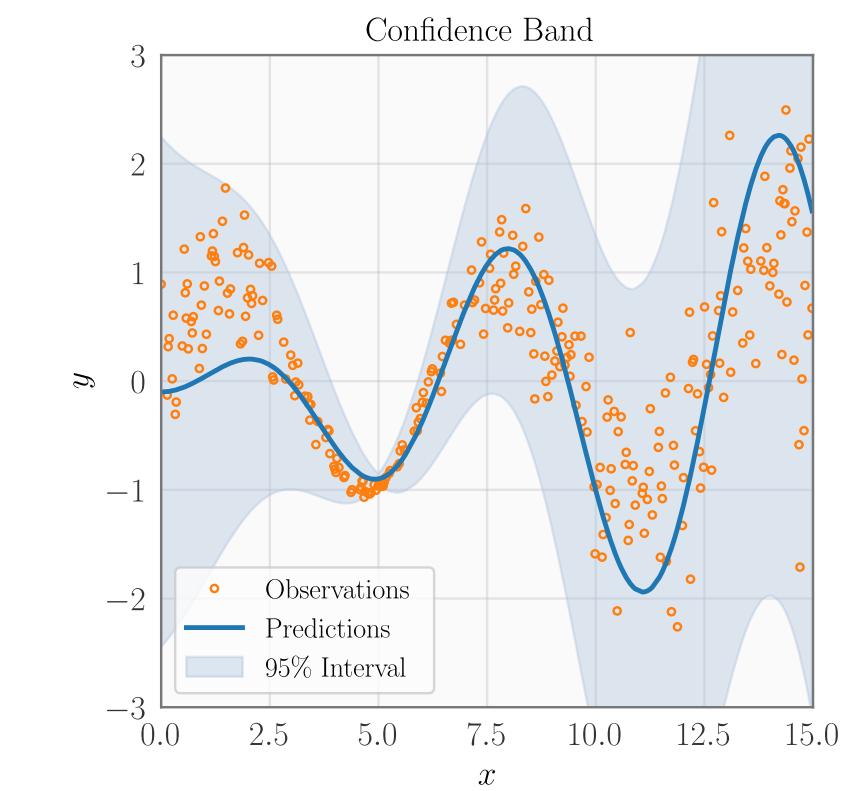
**Laplace model:**  $p(\mathcal{D} | \omega) = \prod_{i=1}^N \frac{1}{2b} e^{-\frac{|y_i - \phi_\omega(x_i)|}{b}} \implies -\nabla_\omega \log p(\mathcal{D} | \omega) \propto \nabla_\omega \text{MAE}(y, \phi_\omega(x))$

# HOW TO MEASURE UNCERTAINTY?

*Selecting the best model is nuanced: one may be overconfident but accurate, while another may be conservative yet inaccurate*



Model A



Model B



# A (POSSIBLE) RECIPE TO MEASURE UNCERTAINTY

## METRIC

*Probabilistic  
Scoring Rules*

*Calibration*

*Correlation*

## WHEN TO USE

*Does the predicted distribution  
match observations?*

*Do confidence levels match  
confidence in reality?*

*Do your uncertainties  
reflect actual mistakes?*

## EXAMPLE

*NLL, CRPS*

*ECE, Miscalibration Area*

*Pearson, Spearman*

# PROBABILISTIC SCORING RULES

A **scoring rule** is a way to measure how good a probabilistic forecast is.

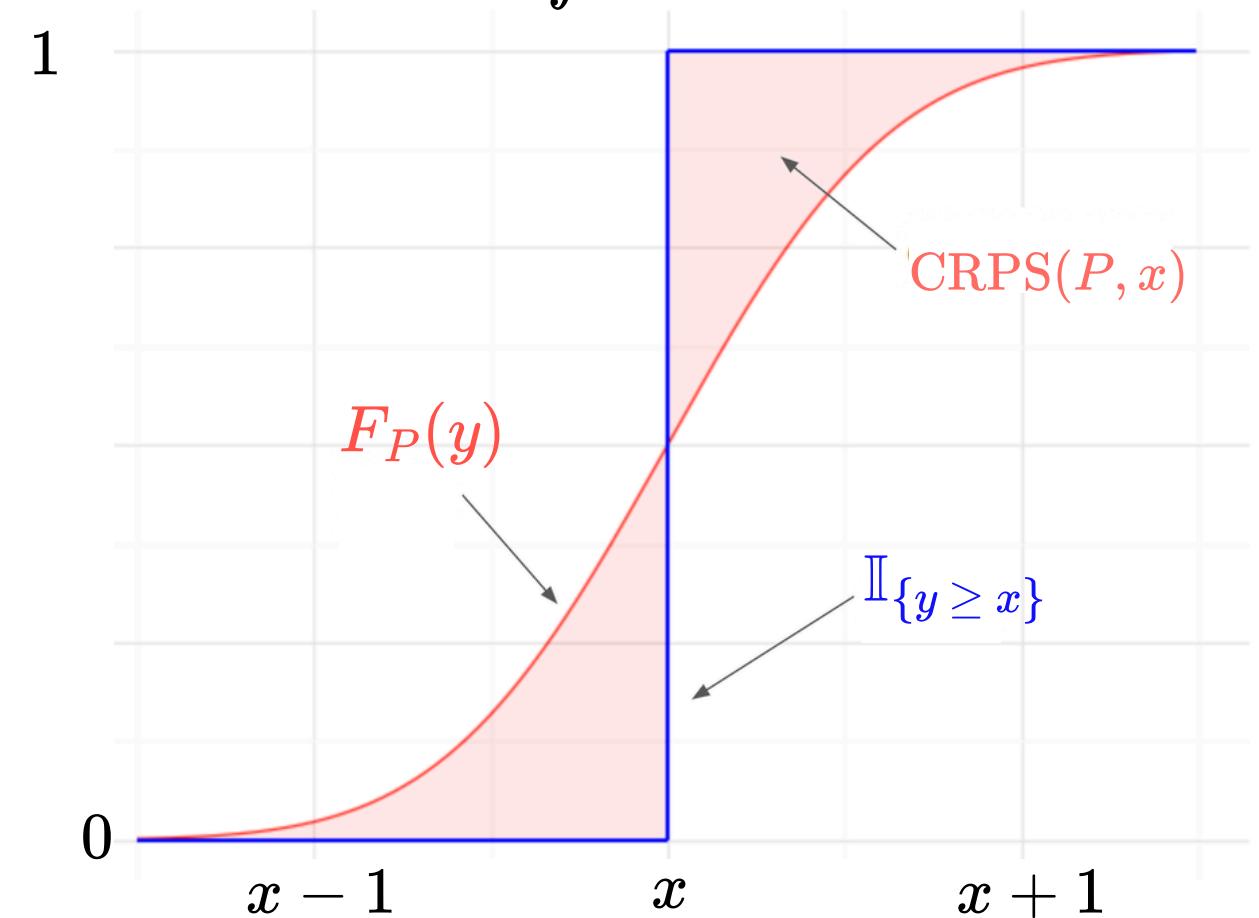
A scoring rule is **proper** if the expected score is maximized when the forecasted probabilities match the forecaster's true beliefs.

**Gaussian Example:**  $P = \mathcal{N}(\mu, \sigma)$

$$\text{NLL}(P, x) = -\log P(x) = \frac{1}{2} \log(2\pi\sigma^2) + \frac{(x - \mu)^2}{2\sigma^2}$$

$$\text{CRPS}(P, x) = \int (F_P(y) - \mathbb{I}_{\{y \geq x\}})^2 dy = \sigma \left[ \frac{1}{\sqrt{\pi}} - 2\phi\left(\frac{x - \mu}{\sigma}\right) - \frac{x - \mu}{\sigma} \left( 2\Phi\left(\frac{x - \mu}{\sigma}\right) - 1 \right) \right]$$

$$\text{CRPS}(P, x) = \int (F_P(y) - \mathbb{I}_{\{y \geq x\}})^2 dy$$



# CALIBRATION

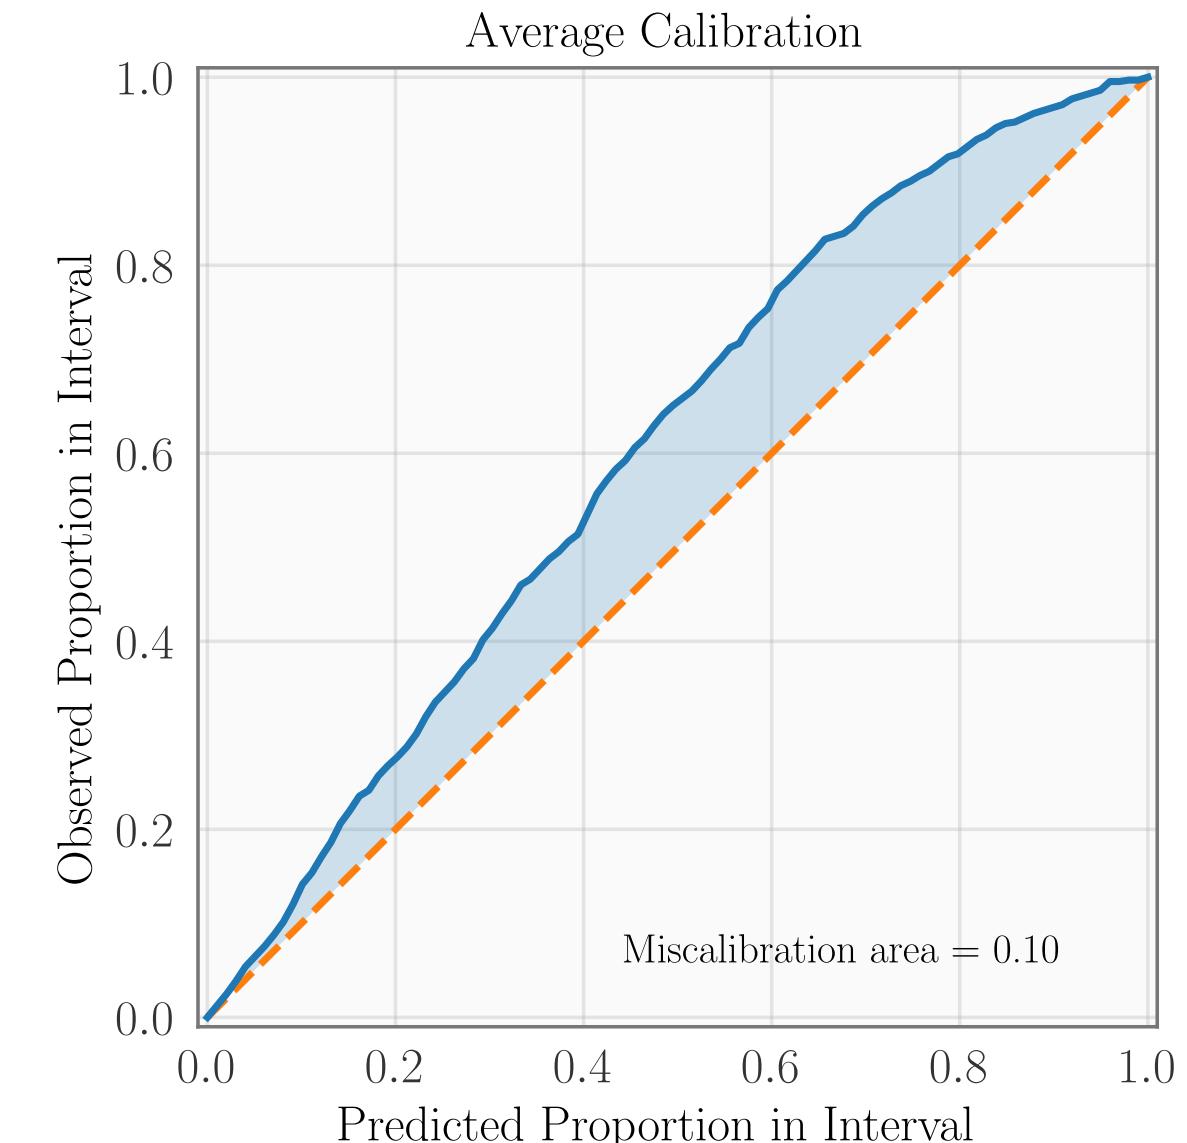
A **calibration** measure tells us how well predicted probabilities align with actual outcomes.

A model is **well-calibrated** if events predicted to occur with probability  $p$  actually happen approximately with proportion  $p$ .

**Gaussian Example:**  $P = \mathcal{N}(\mu, \sigma)$

$$\text{ECE}(P) = \mathbb{E}_{p \sim U(0,1)} [|p - \mathbb{P}(x \leq F_P^{-1}(p))|] = \mathbb{E}_{p \sim U(0,1)} [|p - \mathbb{P}(x \leq \mu + \sigma\sqrt{2}\Phi^{-1}(2p - 1))|]$$

$$\text{Area}(P) = \int_0^1 |p - \mathbb{P}(x \leq F_P^{-1}(p))| dp = \int_0^1 |p - \mathbb{P}(x \leq \mu + \sigma\sqrt{2}\Phi^{-1}(2p - 1))| dp$$

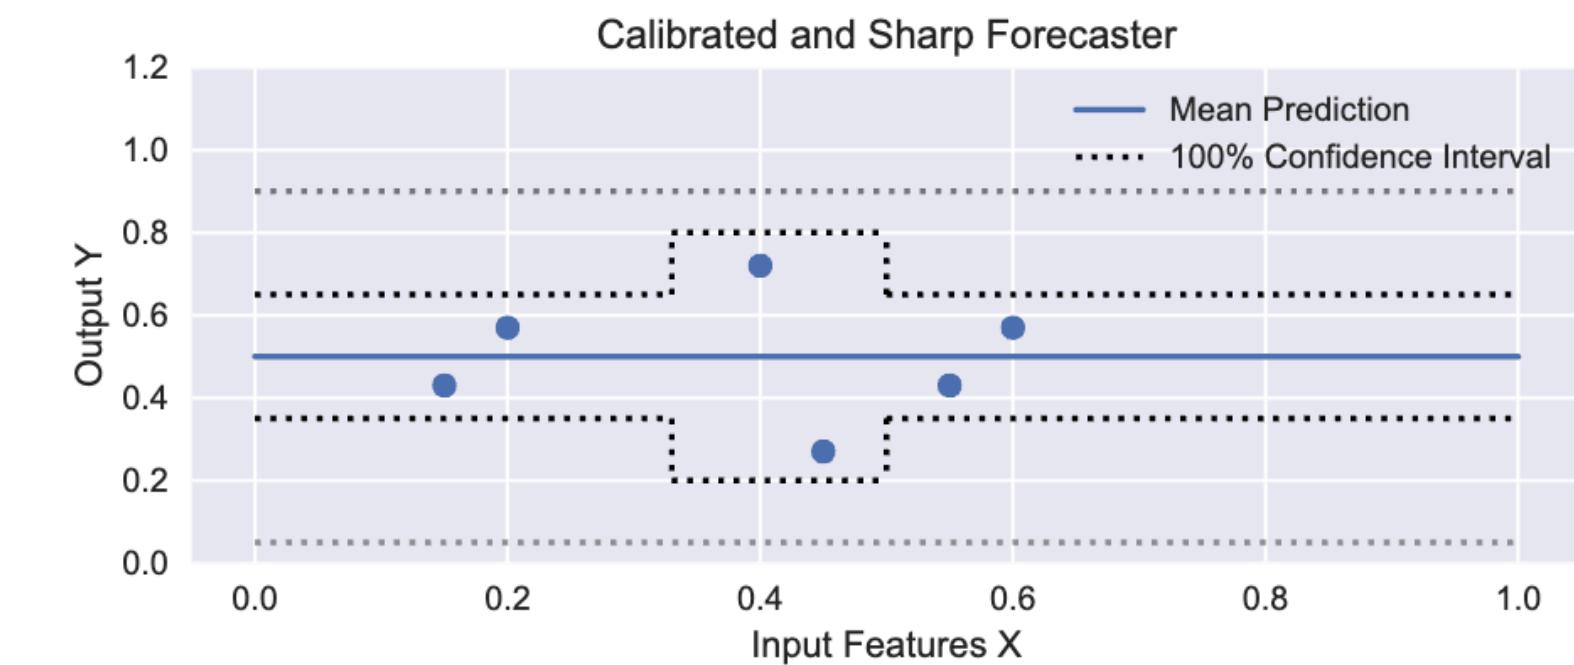
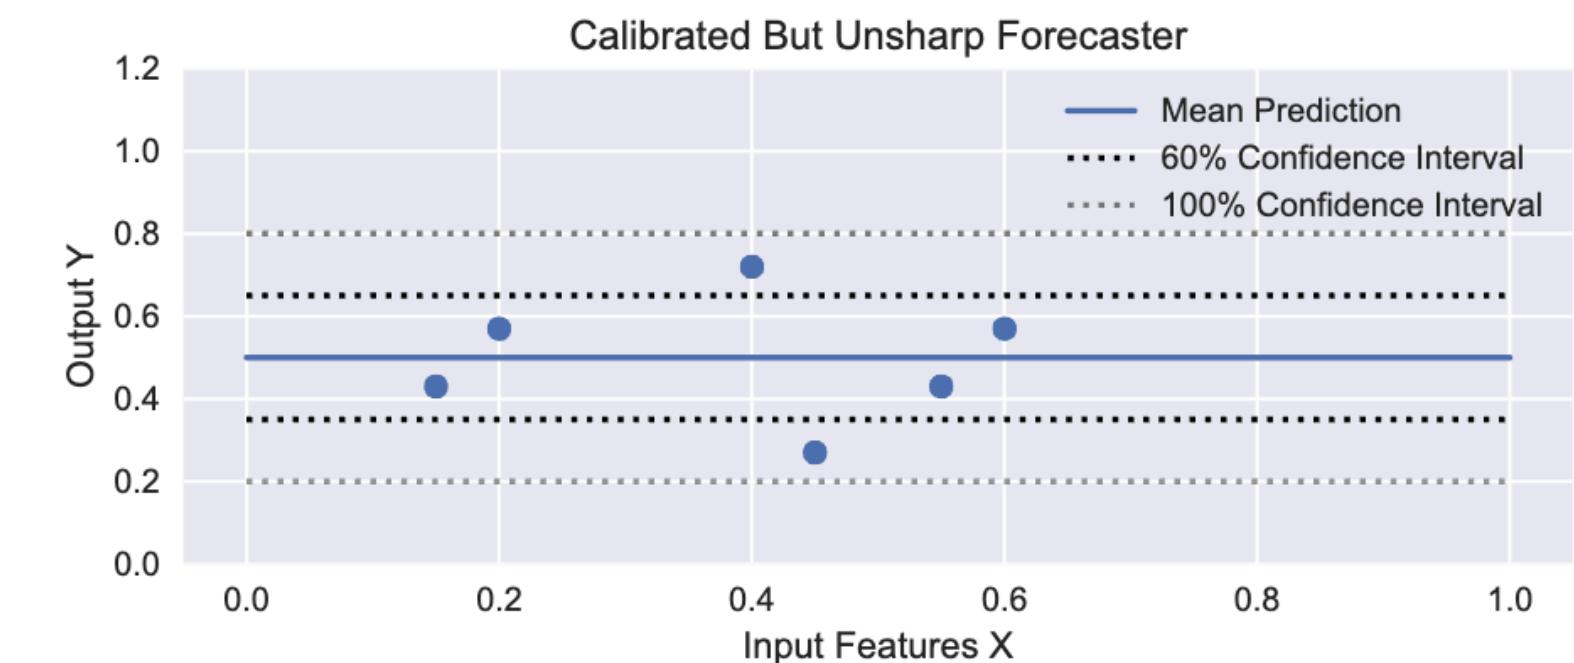


# SHARPNESS

**Calibration by itself is not sufficient to produce a useful forecast!**

**We also want the prediction to be sharp,** meaning it concentrates probability mass tightly around the output.

$$\text{Sharp}(P, x) = \text{Var}_P(x)$$



# CORRELATION

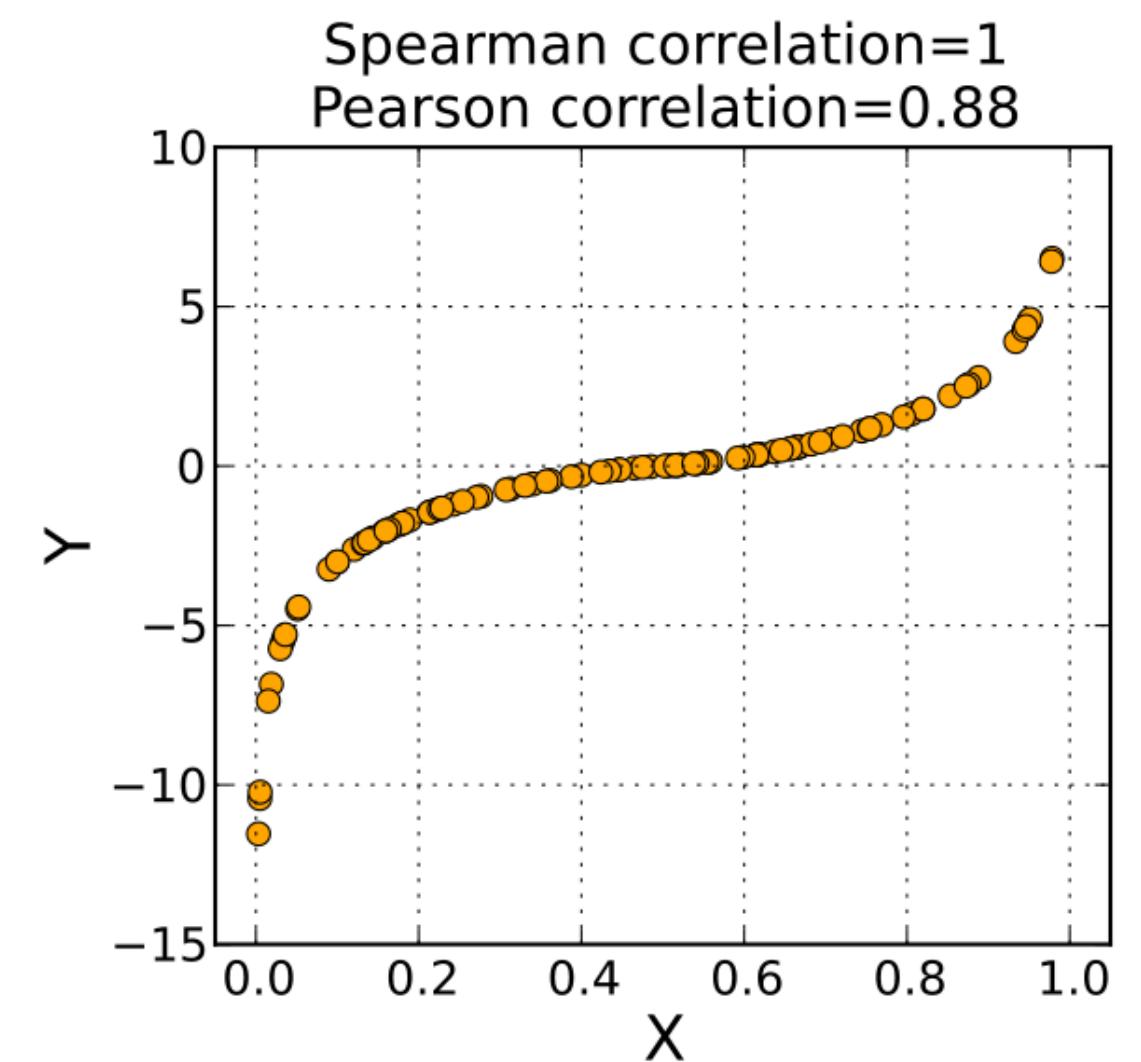
A **correlation coefficient** measures how strongly two variables are related.

The **Pearson correlation** measures how well a straight line describes the relationship between two numerical variables – it's about linear association.

$$\rho_{\text{pearson}} = \frac{\text{Cov}(X, Y)}{\sigma_X \sigma_Y}$$

The **Spearman correlation** measures how well the order of the values in one variable matches the order in another – it's about consistent ranking, even if the relationship is not linear.

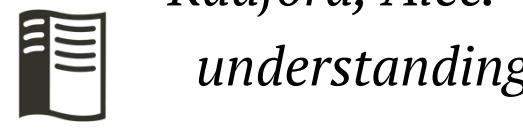
$$\rho_{\text{spearman}} = \frac{\text{Cov}(R[X], R[Y])}{\sigma_{R[X]} \sigma_{R[Y]}}$$



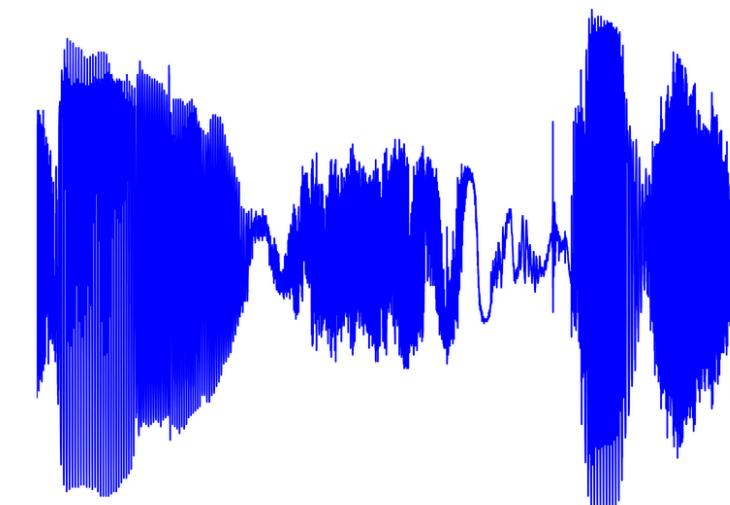
# UNCERTAINTY QUANTIFICATION IN SEQUENCE MODELS



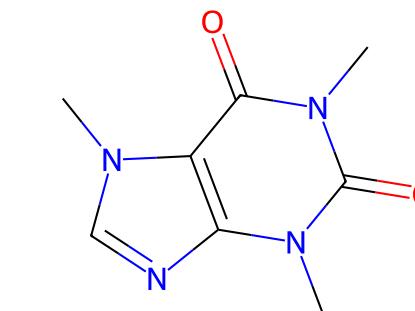
## ChatGPT



*Radford, Alec. "Improving language understanding by generative pre-training."*



*Van Den Oord, Aaron, et al. "Wavenet: A generative model for raw audio."*

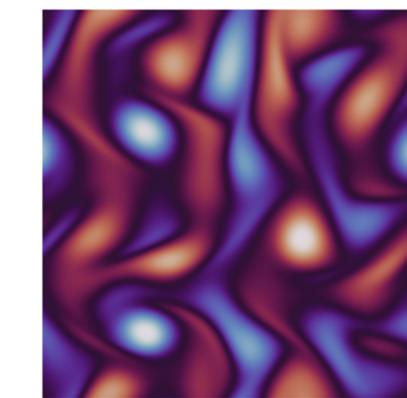


CN1C=NC2=C1C(=O)N(C(=O)N2C)C



*Segler, Marwin HS, et al. "Generating focused molecule libraries for drug discovery with recurrent neural networks."*

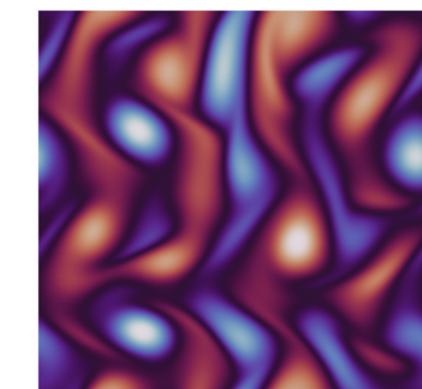
$y_t$



$\mathcal{M}(y_t; \omega)$



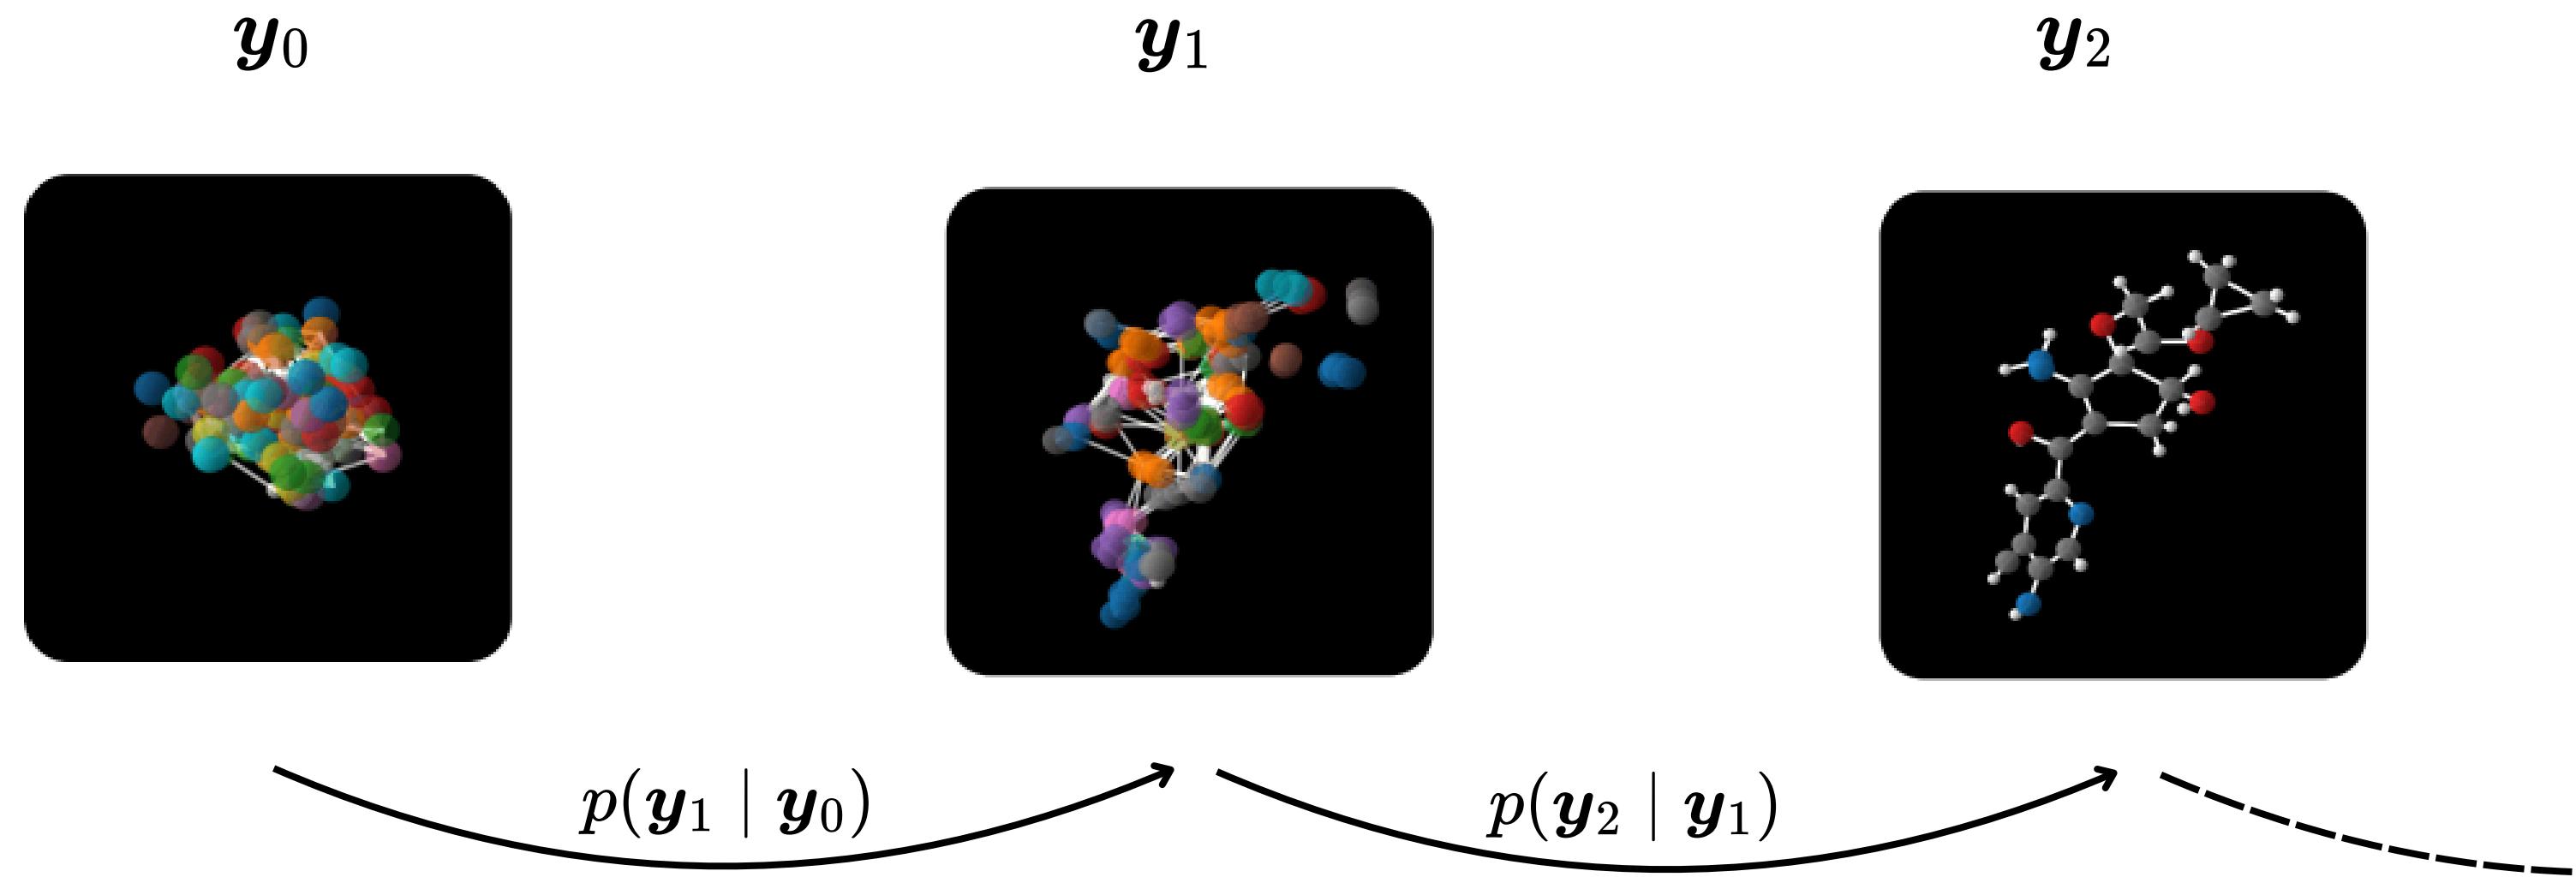
*Brandstetter, Johannes, Daniel Worrall, and Max Welling. "Message passing neural PDE solvers."*



$y_{t+1}$

# AUTOREGRESSIVE AND RECURRENT MODELS

Autoregressive and Recurrent models predict the **next state** in a sequence by processing the **previous states** in that sequence or by keeping a **memory of past states**



# AUTOREGRESSIVE MODELS FOR LANGUAGE

$$p(\mathbf{y}_0, \mathbf{y}_1, \dots, \mathbf{y}_T) = \prod_{t=1}^T p(\mathbf{y}_t \mid \mathbf{y}_{t-1}, \dots, \mathbf{y}_0)$$

The **joint probability** distribution is a **product of factorized distributions** over states

The factorized distribution is usually a **parametric model** with optimizable deterministic weights

The cat is on the  
 $\mathbf{y}_0 \quad \mathbf{y}_1 \quad \mathbf{y}_2 \quad \mathbf{y}_3 \quad \mathbf{y}_4$



table  
 $\mathbf{y}_5$

# AUTOREGRESSIVE MODELS FOR PDES

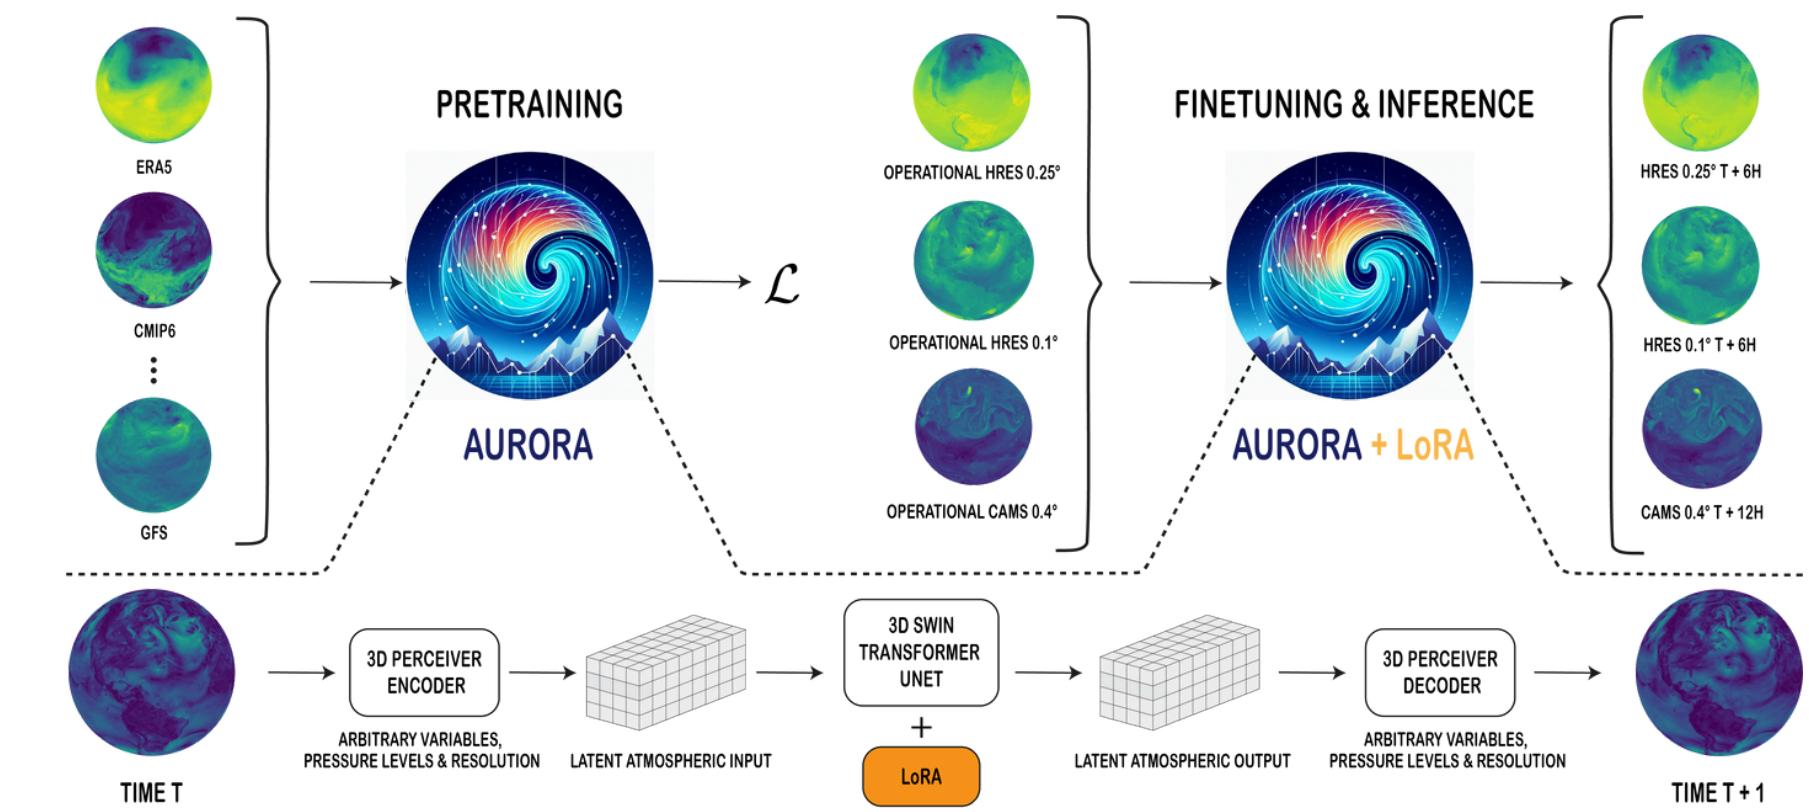
**Autoregressive models** can be used to build Neural Solvers to **solve PDEs** or perform **weather forecast** by iteratively applying the emulator on its output given initial states

Simple Autoregressive Model:

$$p(\mathbf{y}_0, \mathbf{y}_1, \dots, \mathbf{y}_T) = p(\mathbf{y}_0) \prod_{t=2}^T p(\mathbf{y}_t \mid \mathbf{y}_{t-1})$$

where the transition probability is a Neural Network (e.g Transformer based architectures)

$$p(\mathbf{y}_t \mid \mathbf{y}_{t-1}, \mathbf{y}_{t-2}) = \text{NN}(\mathbf{y}_{t-1}, \mathbf{y}_{t-2}; \boldsymbol{\omega})$$



Bodnar, Cristian, et al. "A foundation model for the Earth system." *Nature*, 2025

# PDE-REFINER

## Refinement Process

The multistep refinement process allows more accurate modelling of all frequency components

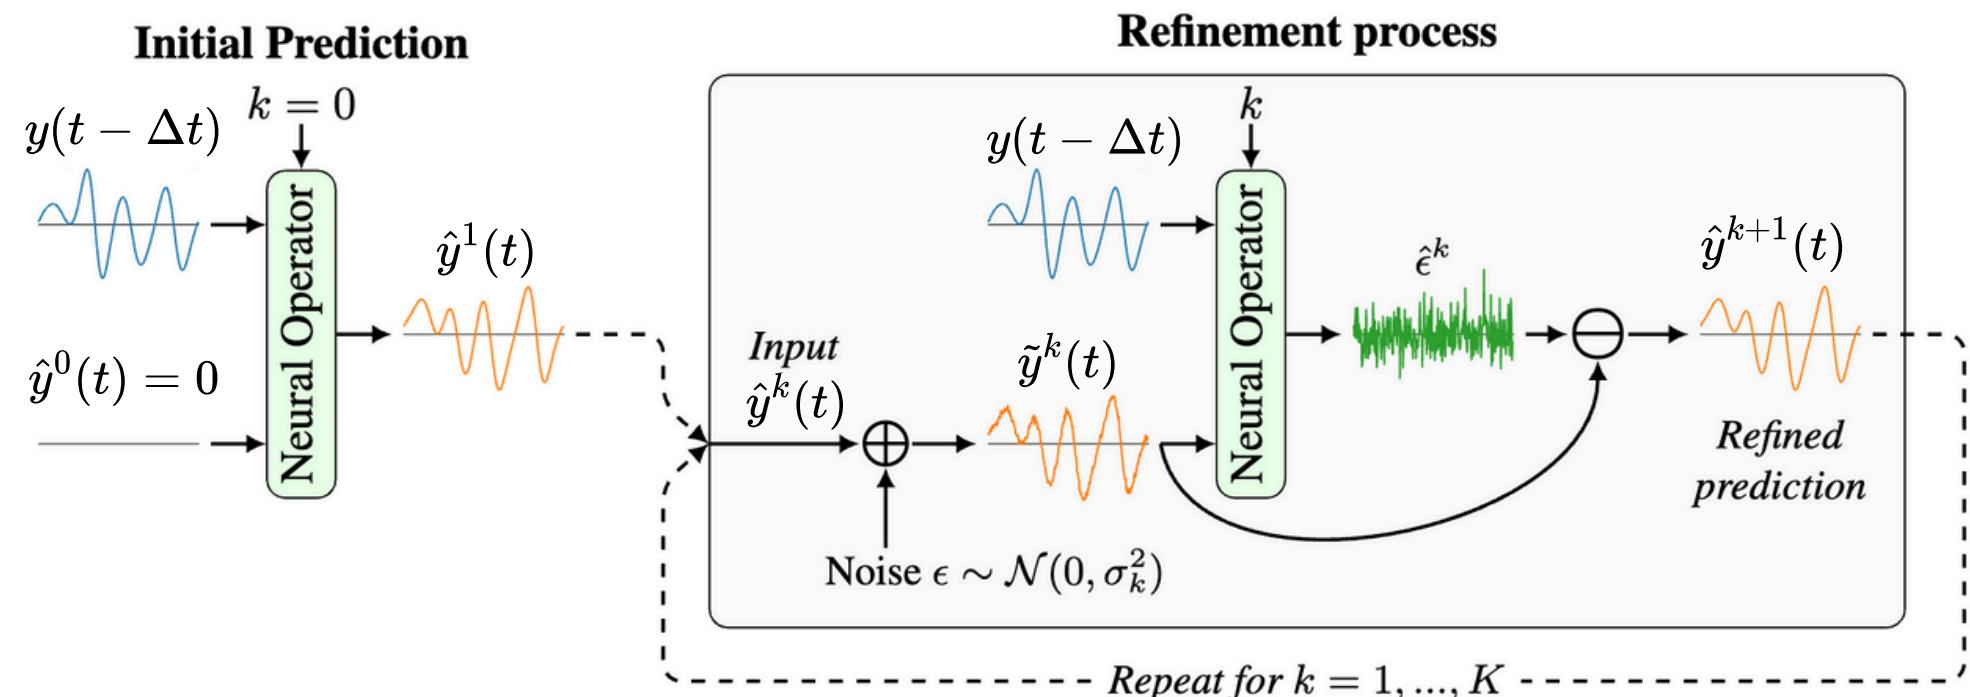
## Denoising Loss Function

$$\mathcal{L}(\theta) = \mathbb{E}_{k \sim U(0, K)} [\mathcal{L}_k(\theta)]$$

with,

$$\mathcal{L}_0(\theta) = \|y_t - \text{NN}(\hat{y}^0(t), y(t - \Delta t), 0)\|^2 \quad (\text{one-step loss})$$

$$\mathcal{L}_k(\theta) = \|\epsilon_k - \text{NN}(y(t) + \sigma_k \epsilon, y(t - \Delta t), k)\|^2, k = 1, \dots, K \quad (\text{refinement loss})$$



Lippe, Phillip, et al. "PDE-Refiner: Achieving Accurate Long Rollouts with Neural PDE Solvers", 2023.

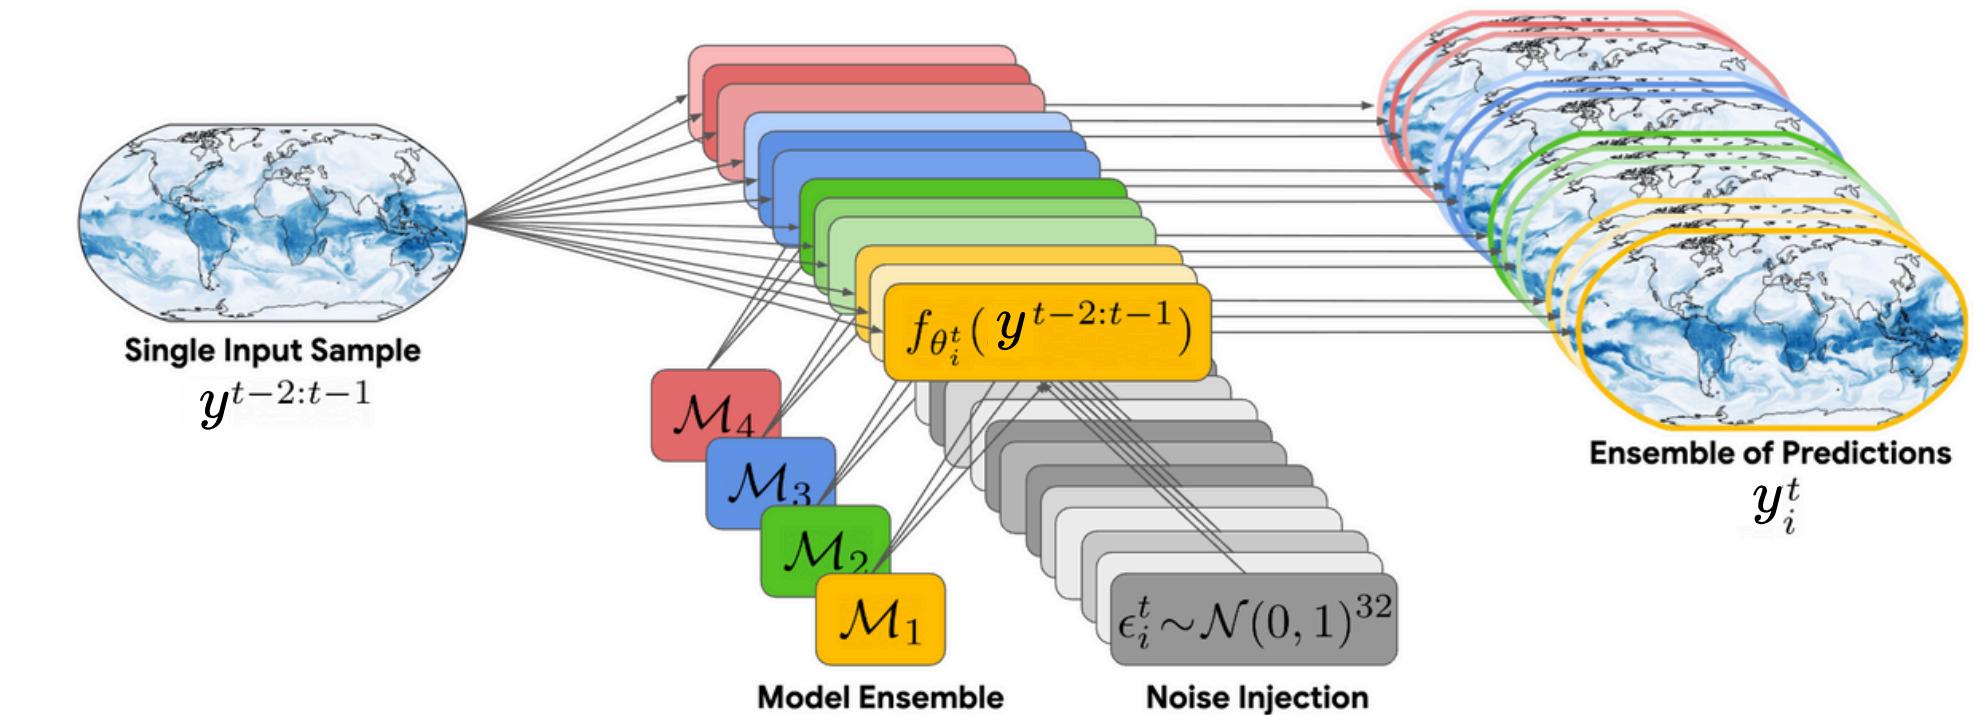
# FUNCTIONAL GENERATIVE NETWORKS

## Ensemble Uncertainty

$$p(y^t | y^{t-2:t-1}) = \sum_i p(y^t | y^{t-2:t-1}, \omega) \delta(\omega - \theta_i)$$

## Perturbation Uncertainty (only LayerNorm)

$$\theta_i^t = \theta_i + \Delta_i \cdot \epsilon_i^t, \quad \epsilon_i^t \sim \mathcal{N}(0, 1)$$



Alet, Ferran, et al. "Skillful joint probabilistic weather forecasting from marginals", 2025.

## Trained to minimize CRPS

$$\mathcal{L}(\theta_i, \Delta_i) = \frac{1}{N} \sum_n |y_{n;(\theta_i, \Delta_i)}^t - y^t| - \frac{1}{2N(N-1)} \sum_{n,n'} |y_{n;(\theta_i, \Delta_i)}^t - y_{n';(\theta_i, \Delta_i)}^t|$$

# BARNN

## Joint Weight-States update

Jointly sample in an alternating way  
weights and states

$$p(\mathbf{y}_0, \omega_1, \mathbf{y}_1, \omega_2, \dots, \mathbf{y}_T) = \prod_{t=1}^T p(\mathbf{y}_t | \mathbf{y}_{<t}, \omega_t) p(\omega_t)$$

## Variational Inference

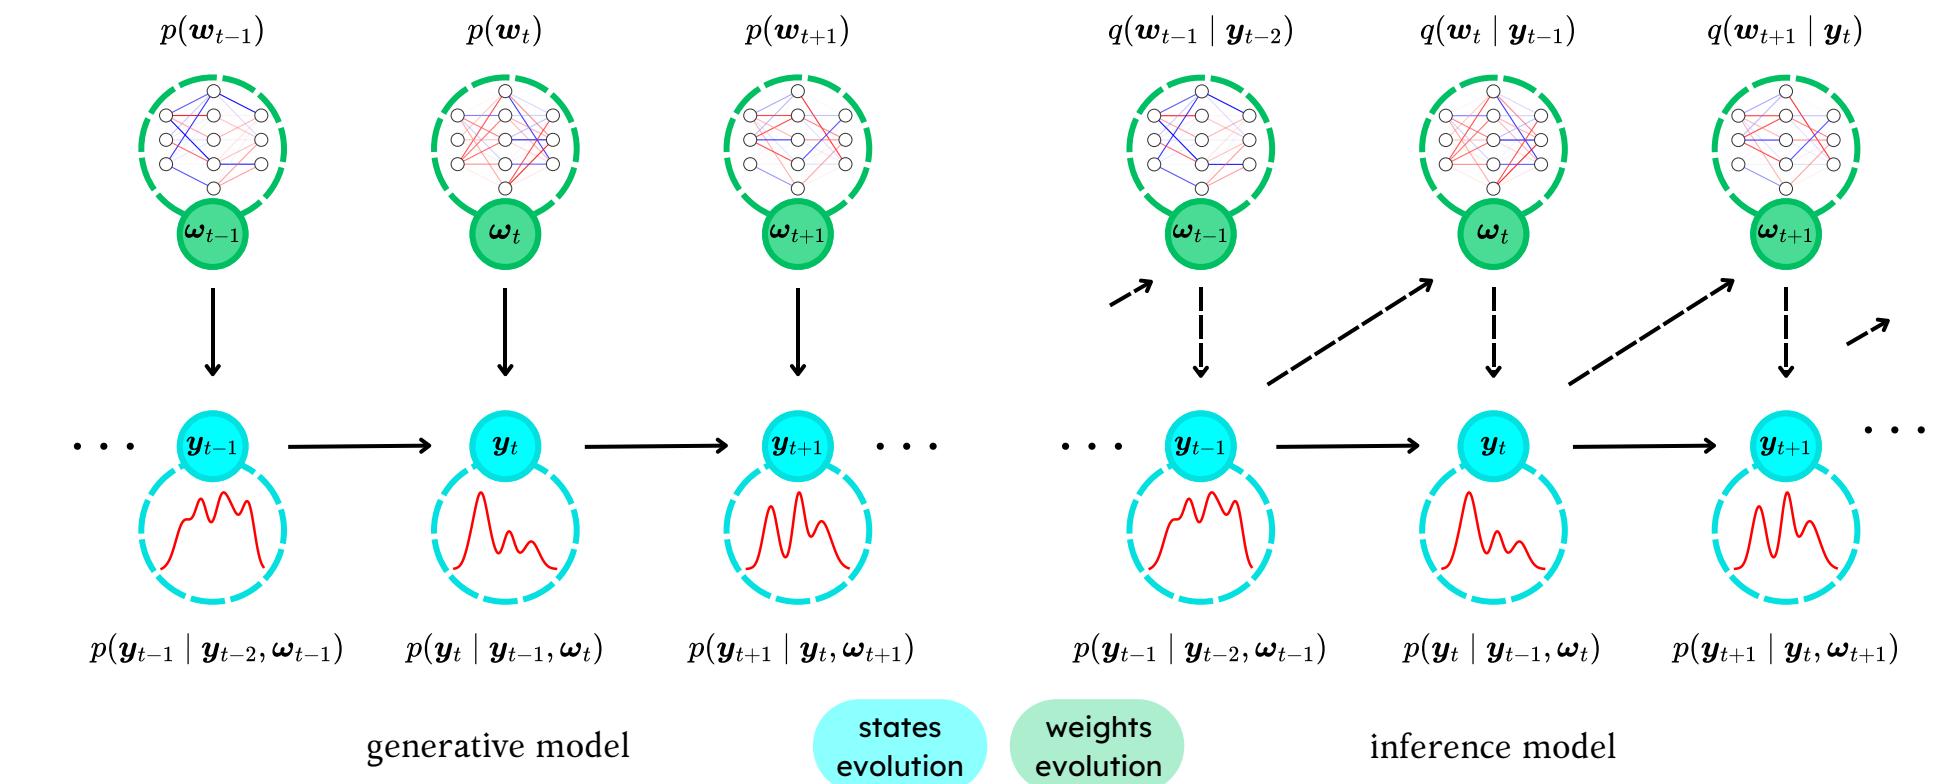
Use variational inference to optimise the  
neural network

$$q_\phi(\omega_1, \dots, \omega_T | \mathbf{y}_{<T}) = \prod_{t=1}^T q(\omega_t | \mathbf{y}_{<t}) \quad \text{and,} \quad p(\omega_t) = \int q(\omega_t | \mathbf{y}_{<t}) p(\mathbf{y}_{<t}) d\mathbf{y}_{<t}$$

$$\mathcal{L}(\phi) = \mathbb{E}_{t \sim \mathcal{U}[1,T]} [\mathbb{E}_{\omega_t \sim q_\phi} [\log p(\mathbf{y}_t | \mathbf{y}_{0:t}, \omega_t)] - D_{KL} [q_\phi(\omega_t | \mathbf{y}_{<t}) \| p(\omega_t)]]$$

## Scale to Large Networks

Local-reparametrization allos model's scale to large Neural Network  $q(\omega_t^l | \mathbf{y}_{<t}) = \mathcal{N}(\alpha_t^l \Omega^l, (\alpha_t^l \Omega^l)^2)$   $\alpha_t^l = f(\mathbf{y}_{<t}; \delta)$

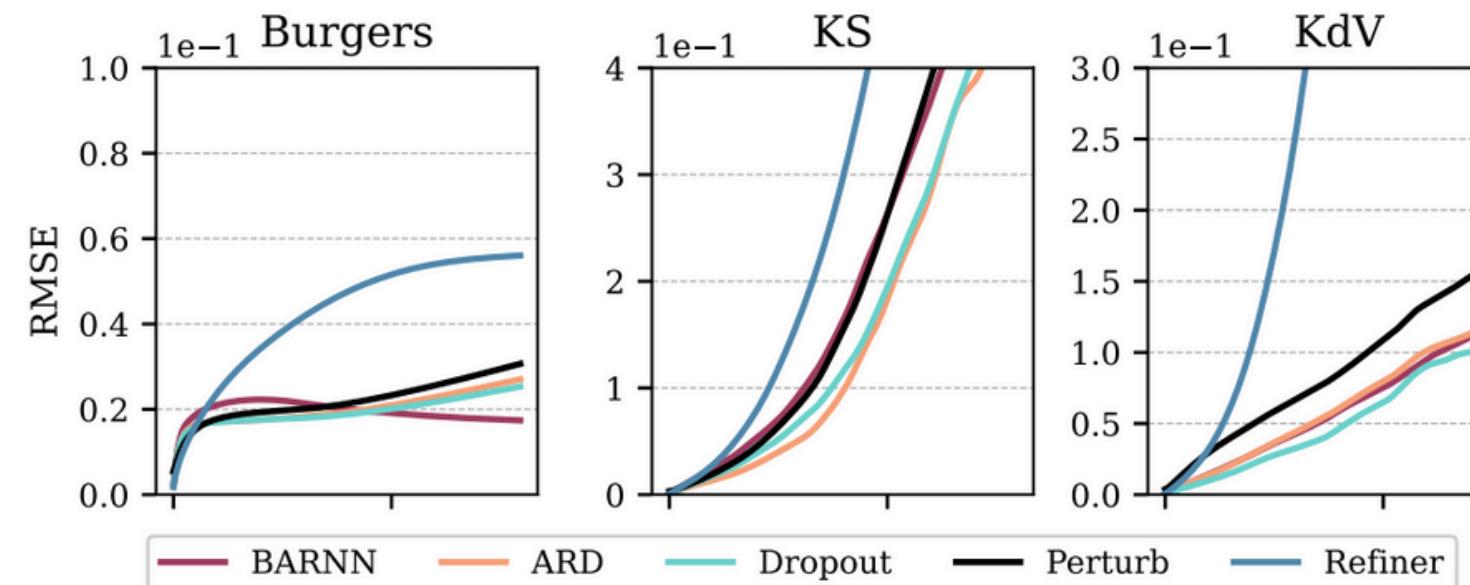


*Coscia, Dario, et al. "BARNN: A Bayesian Autoregressive and Recurrent Neural Network", 2025.*

# UNCERTAINTY ON PDE MODELLING

Most models can **accurately solve PDEs** with **low point-wise RMSE** but **struggles in UQ scenarios** where lower NLL and ECE is expected to have reliable uncertainties

<i>Burgers</i>	$\partial_t u + u \partial_x u - \nu \partial_{xx} u = 0$
<i>Kuramoto Sivashinsky</i>	$\partial_t u + u \partial_x u + \partial_{xx} u + \partial_{xxxx} u = 0$
<i>Korteweg de Vries</i>	$\partial_t u + u \partial_x u + \partial_{xxx} u = 0$



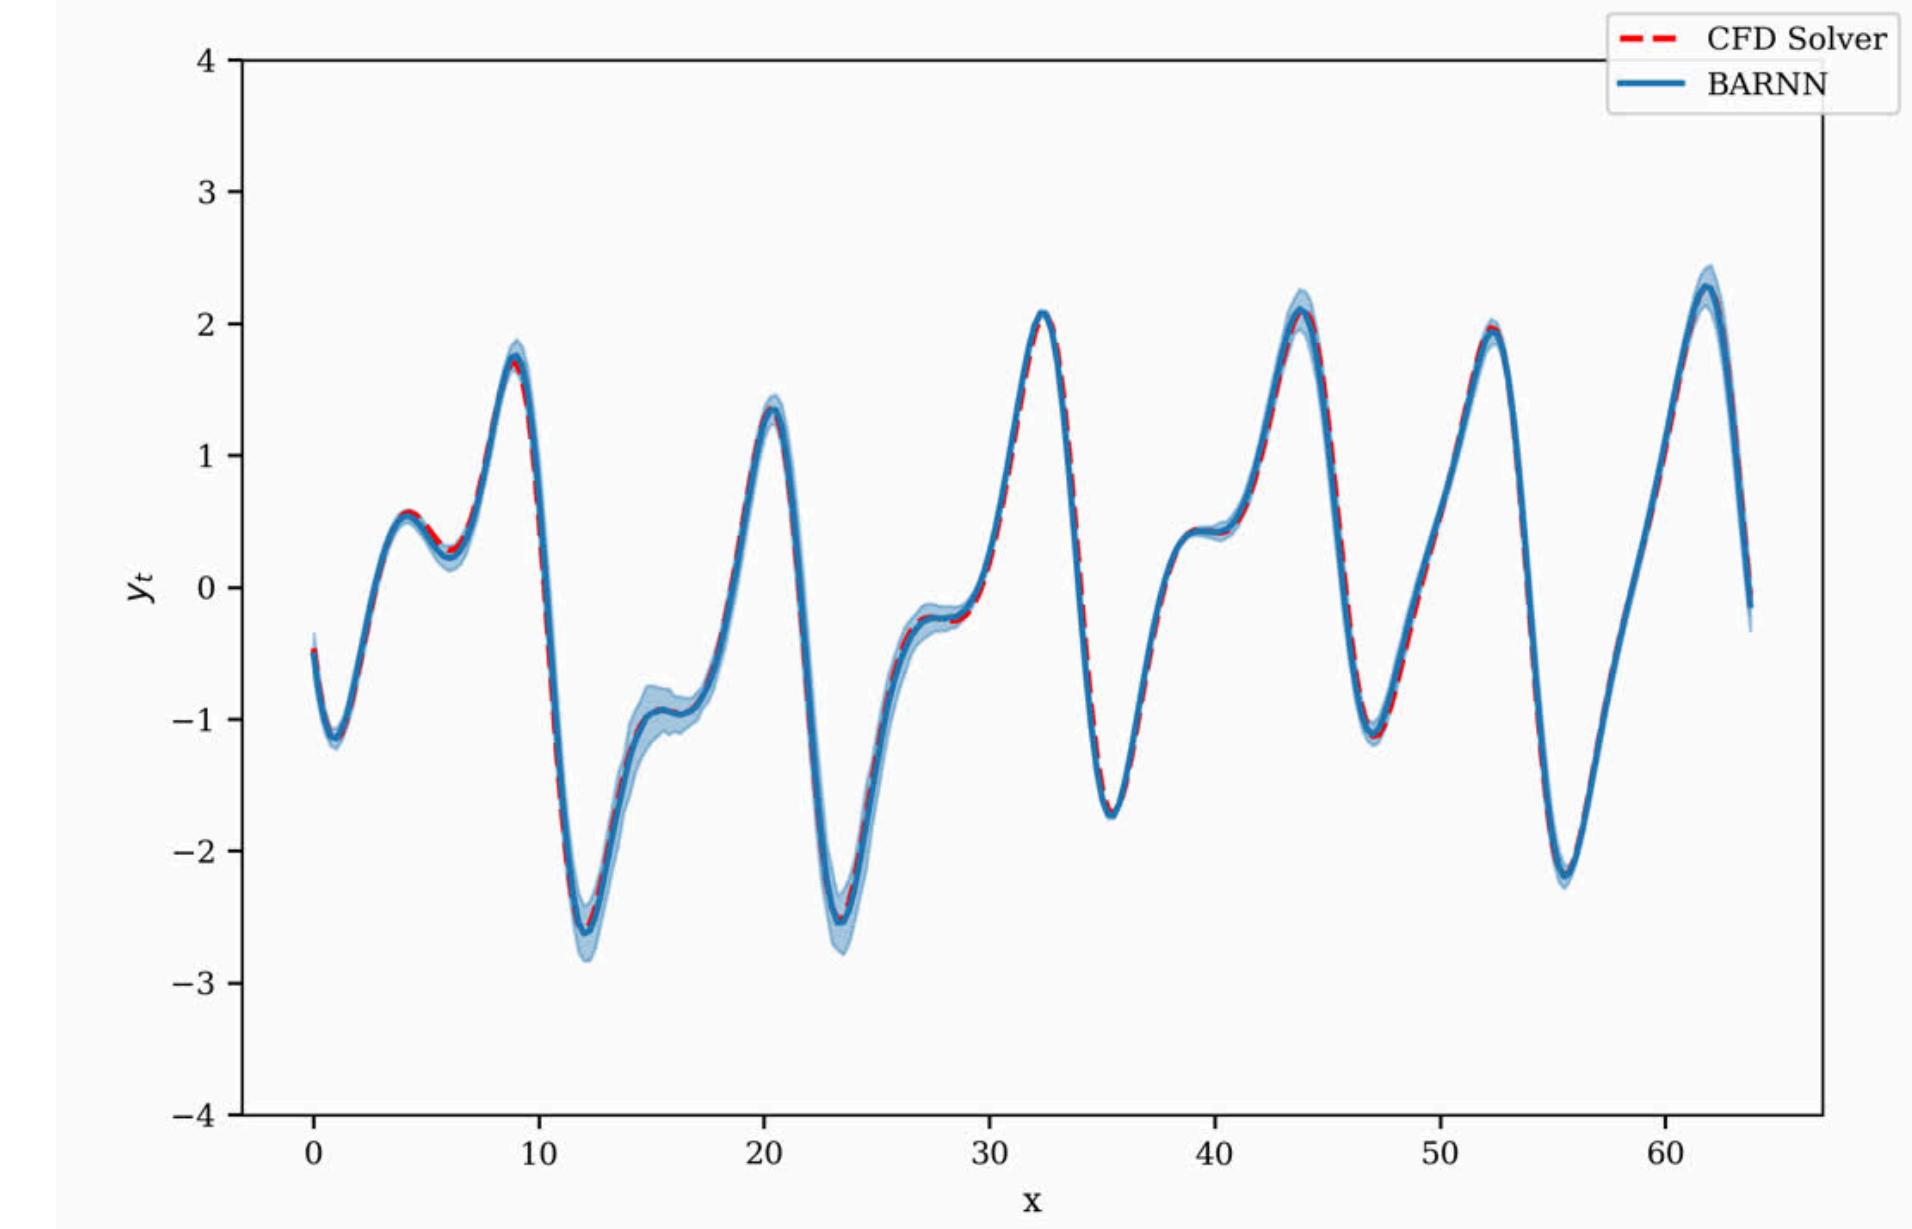
Model	NLL ( $\downarrow$ )			ECE ( $\downarrow$ )		
	Burgers	KS	KdV	Burgers	KS	KdV
BARNN	<b><math>-2.51^{\pm 0.29}</math></b>	<b><math>-0.87^{\pm 0.52}</math></b>	<b><math>-0.94^{\pm 1.23}</math></b>	<b><math>0.05^{\pm 0.03}</math></b>	<b><math>0.04^{\pm 0.03}</math></b>	<b><math>0.10^{\pm 0.01}</math></b>
ARD	$5.00^{\pm 7.11}$	$10.7^{\pm 11.6}$	break	$0.26^{\pm 0.10}$	$0.17^{\pm 0.01}$	$0.19^{\pm 0.05}$
Dropout	$0.06^{\pm 4.59}$	$-0.48^{\pm 0.65}$	$0.40^{\pm 4.57}$	$0.17^{\pm 0.15}$	$0.10^{\pm 0.01}$	$0.14^{\pm 0.07}$
Perturb	$-2.01^{\pm 0.08}$	$-0.65^{\pm 1.07}$	$-0.37^{\pm 1.32}$	$0.16^{\pm 0.01}$	$0.05^{\pm 0.02}$	$0.11^{\pm 0.06}$
Refiner	break	break	break	$0.30^{\pm 0.08}$	$0.23^{\pm 0.02}$	$0.34^{\pm 0.09}$

# UNCERTAINTY ON PDE MODELLING

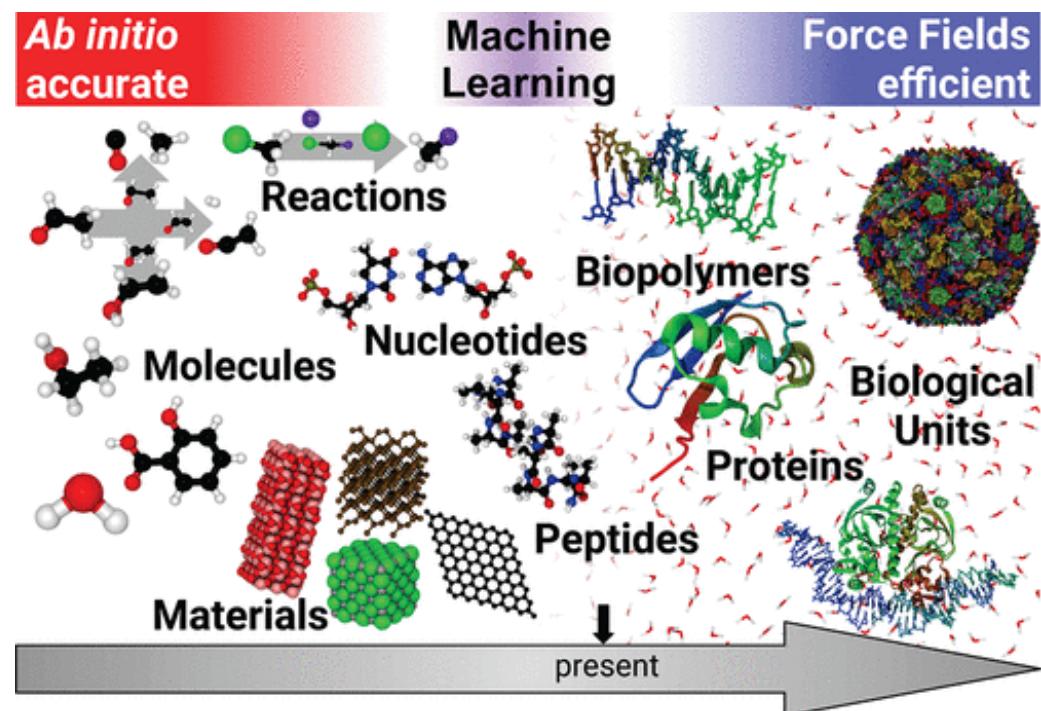
***Confidence intervals are tighter in regions of small uncertainties and larger otherwise***

(using BARNN model)

*Kuramoto Sivashinsky*       $\partial_t u + u \partial_x u + \partial_{xx} u + \partial_{xxxx} u = 0$



# UNCERTAINTY QUANTIFICATION IN MATERIAL SCIENCE



Unke, Oliver T., et al. "Machine learning force fields." *Chemical Reviews* (2021)



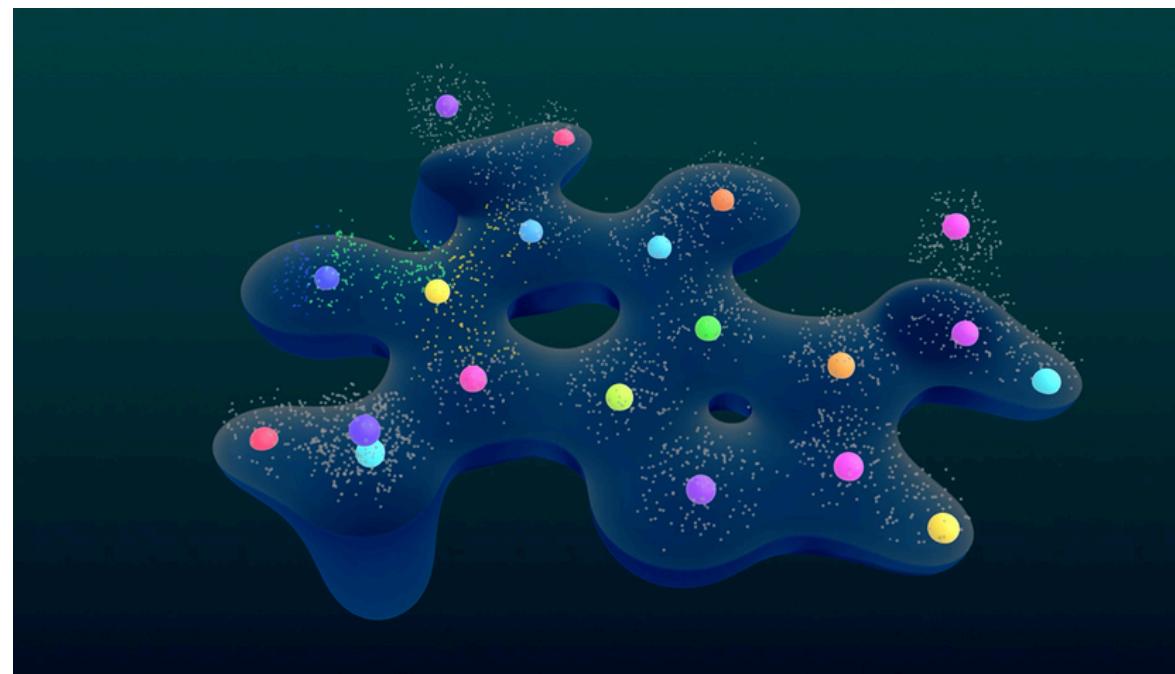
Wood, Brandon M., et al. "UMA: A Family of Universal Models for Atoms." arXiv preprint arXiv:2506.23971 (2025).



Rhodes, Benjamin, et al. "Orb-v3: atomistic simulation at scale." arXiv preprint arXiv:2504.06231 (2025).

# A SMALL INTRO TO MATERIAL SCIENCE

**Atoms form the basic building blocks of molecules and materials, held together by their electrons.** This “electron glue” determines stability and key properties. **Accurately modeling it helps predict chemical reactions, drug effectiveness, material performance for carbon capture, and improvements in renewable energy storage**



Luise, Giulia, et al. "Accurate and scalable exchange-correlation with deep learning."  
arXiv preprint arXiv:2506.14665 (2025).

*Kohn-Sham formalism is used to compute energy of the system*

$$E = \min_{\rho} E_{\text{tot}}[\rho], \quad E_{\text{tot}}[\rho] = \int v(r)\rho(r) dr + \frac{1}{2} \int \int \frac{\rho(r)\rho(r')}{|r - r'|} dr dr' + T_s[\rho] + E_{\text{xc}}[\rho]$$

*where,*

$\rho$     electron density

$E_{\text{xc}}$     exchange correlation energy

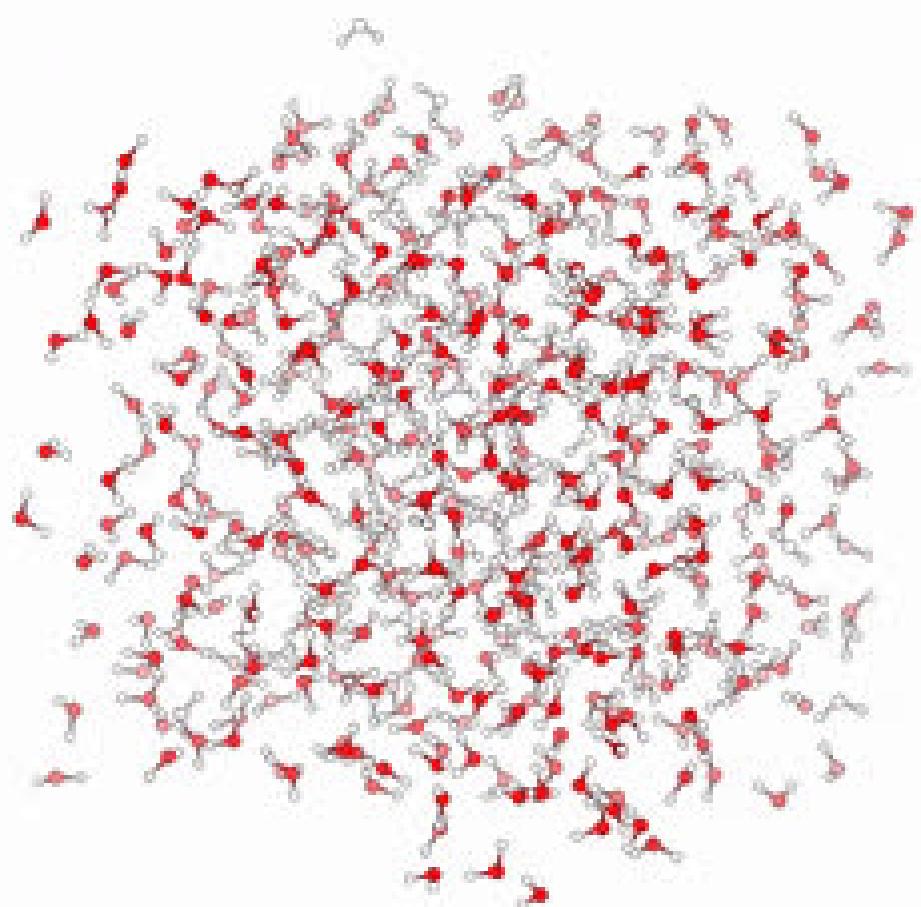
$E_{\text{tot}}$     electron density

$T_s$     kinetic energy of the system

$v(r)$     nuclei external potential

# A SMALL INTRO TO MATERIAL SCIENCE

**Molecular dynamics (MD)** is a simulation method that applies **Newton's laws of motion** using **quantum computed Forces** to track how atoms and molecules move over time. It's like running a "virtual experiment," where you observe how a system of interacting particles evolves to explore their behavior and properties.



A simple MD simulation of water (H2O)

Forces are computed by taking the negative gradient of the Energy (usually computed via DFT)

$$\mathbf{F} = -\nabla_{\mathbf{r}} E_{\text{tot}}$$

Then using Newton's law of motion we study the material over time

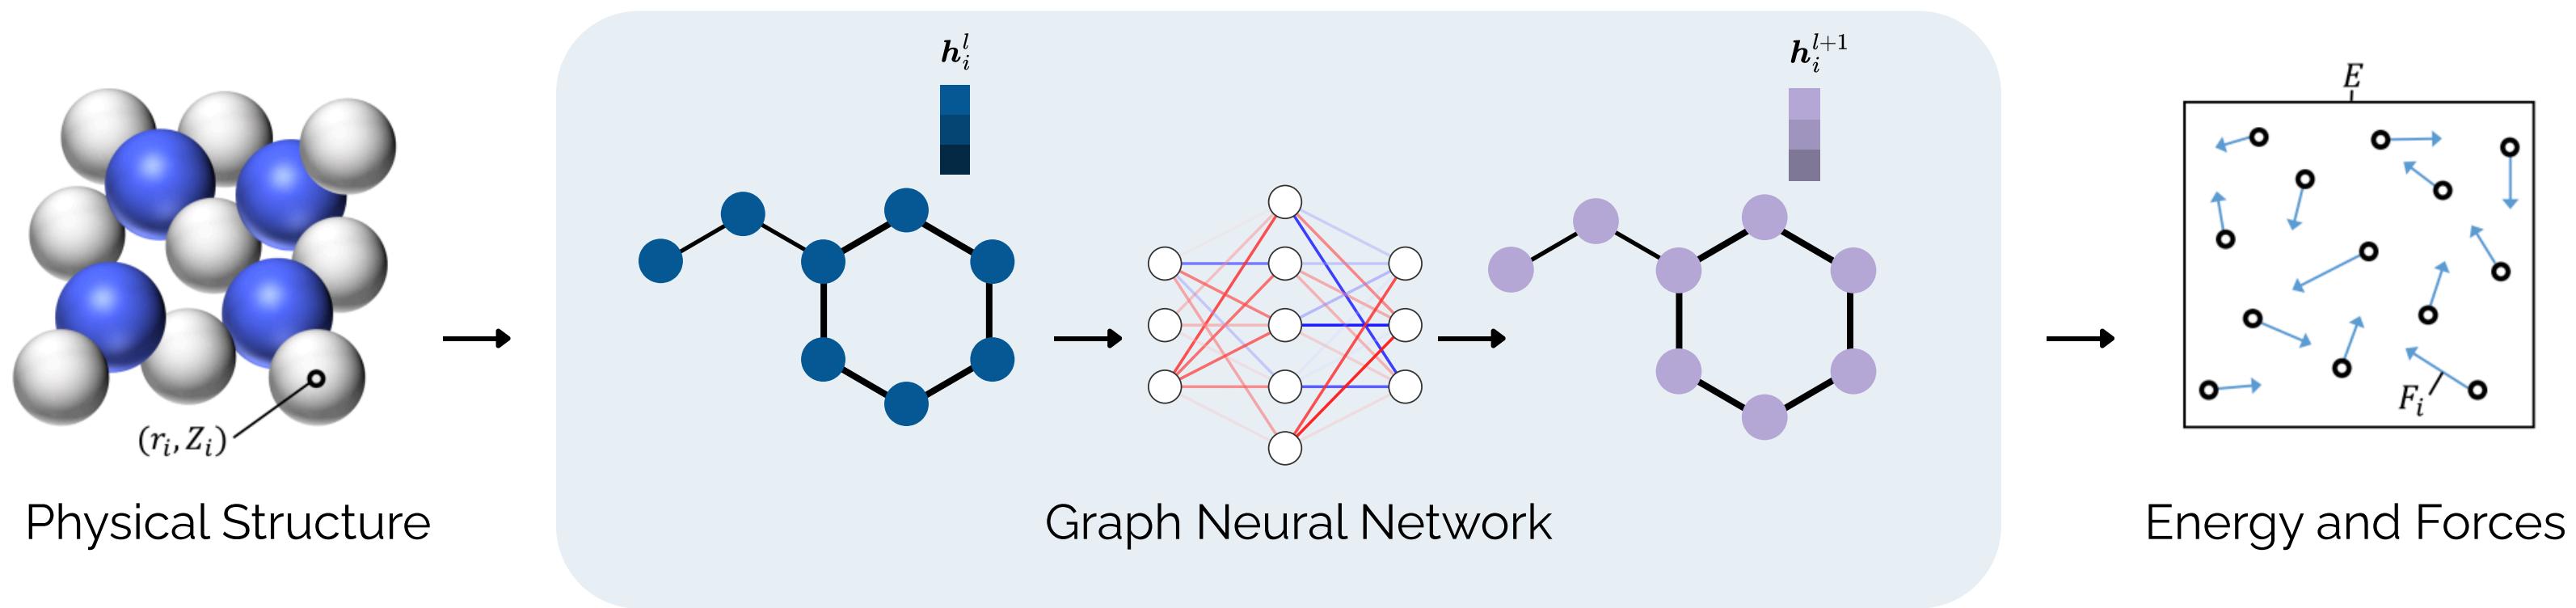
$$\frac{d\mathbf{r}}{dt} = \mathbf{v}, \quad \frac{d\mathbf{v}}{dt} = \frac{\mathbf{F}}{m}$$

# MACHINE LEARNING FORCE FIELDS

**Machine Learning Force Fields** use **data-driven models** to predict **structure energy and forces** with high accuracy and efficiency. There are two types of Models:

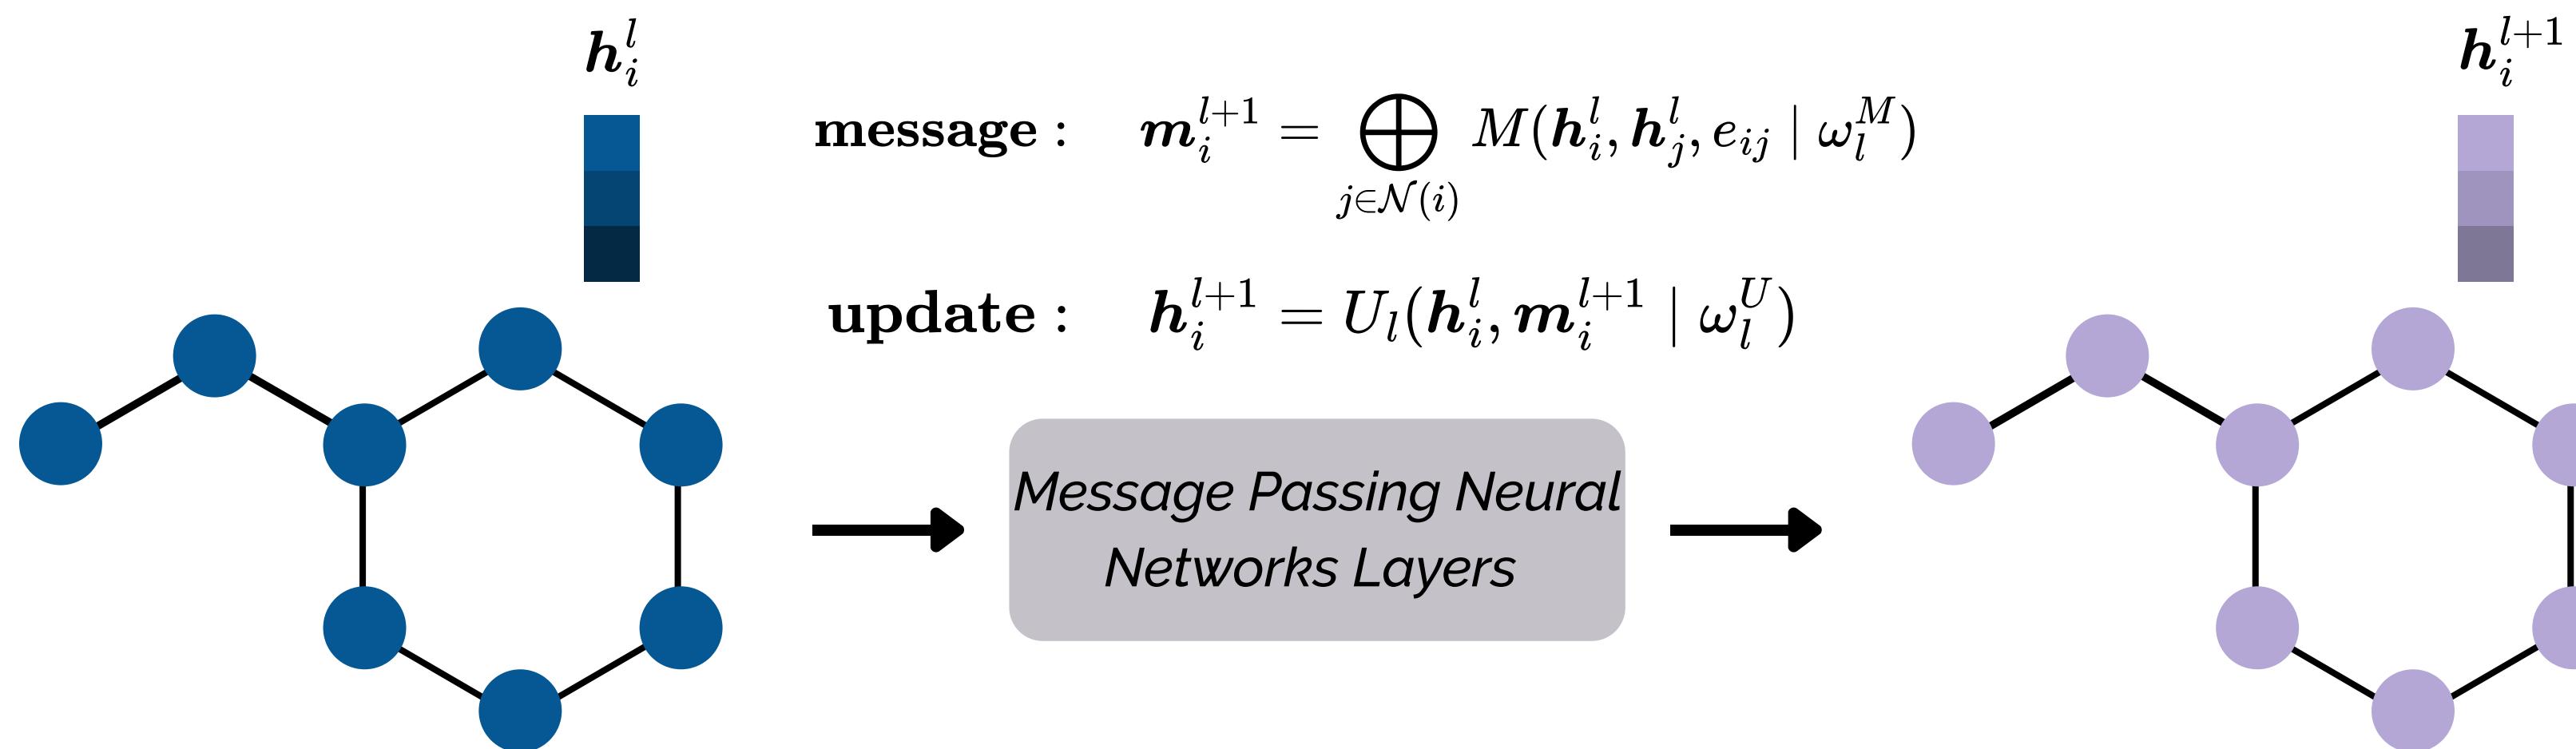
**Conservative Models:** Network predicts Energy. The Forces are obtain by taking the negative gradient with respect to positions using automatic differentiation. Example: **MACE, UMA, ...**

**Direct Models:** Network predicts directly Energy and Forces. Example: **ORB, ...**



# TRAINING MESSAGE PASSING NETWORKS

**Message Passing Neural Networks** (MPNNs) are trained to learn interactions by iteratively exchanging information (messages) between nodes in a graph. They capture complex local environments and many-body effects with high flexibility.



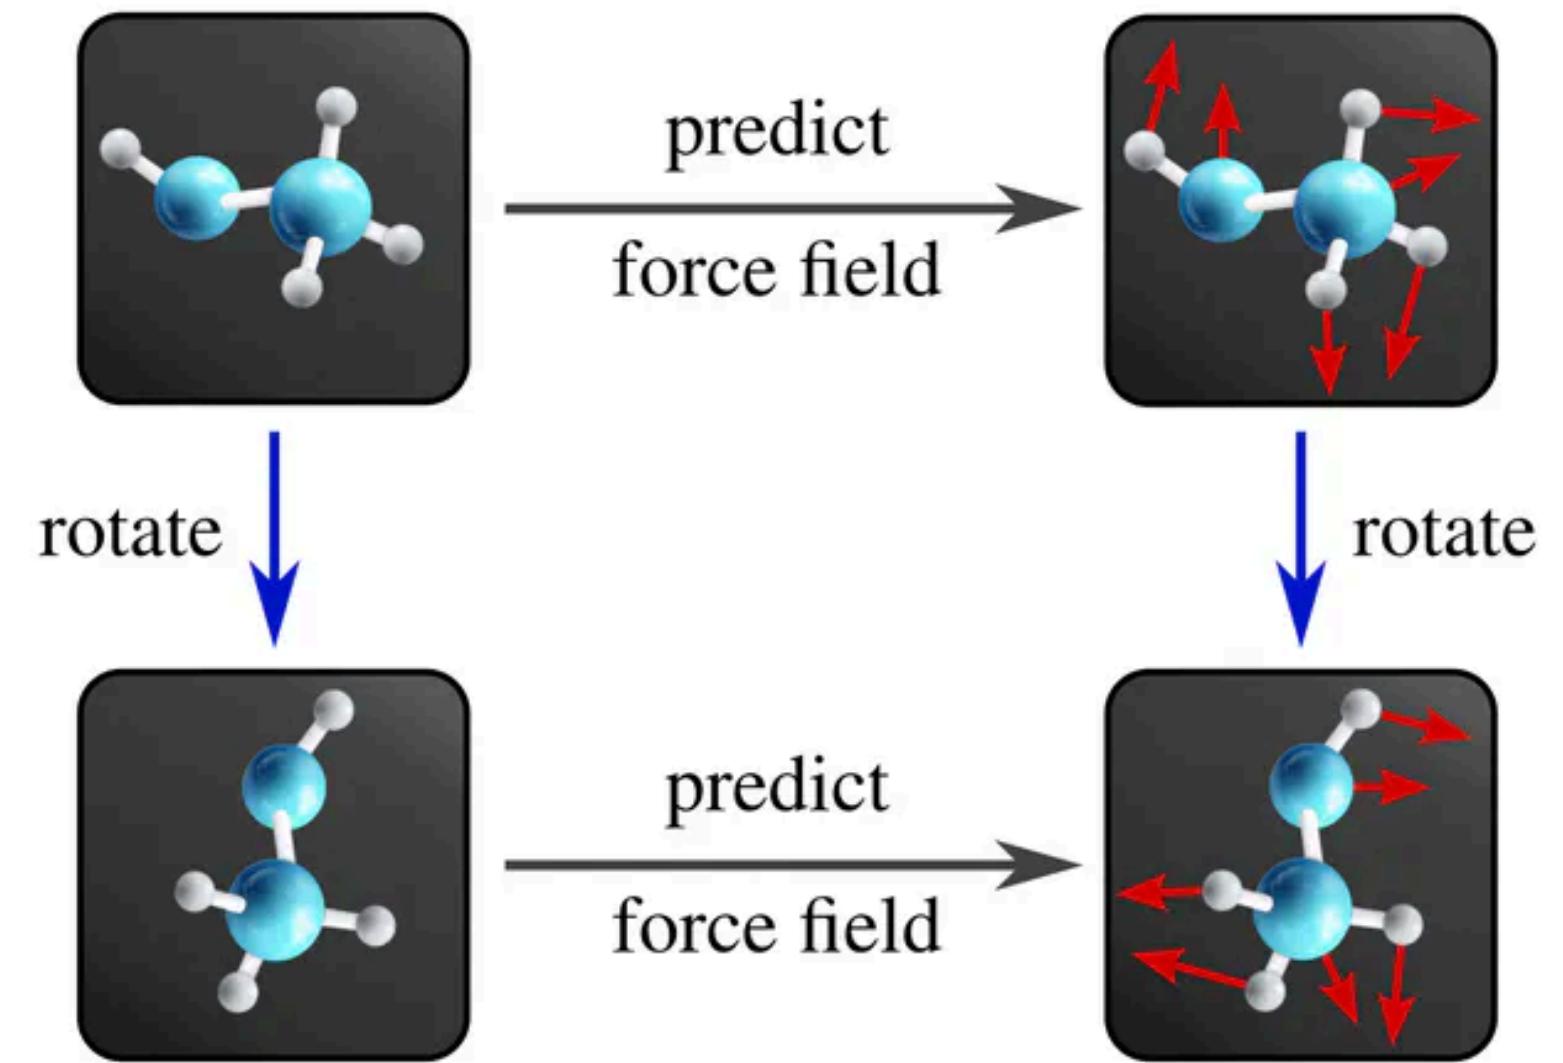
# IS IT ALL?

**Machine Learning Force Fields** are usually built to satisfy SO(3) equivariance/invariance

## Example

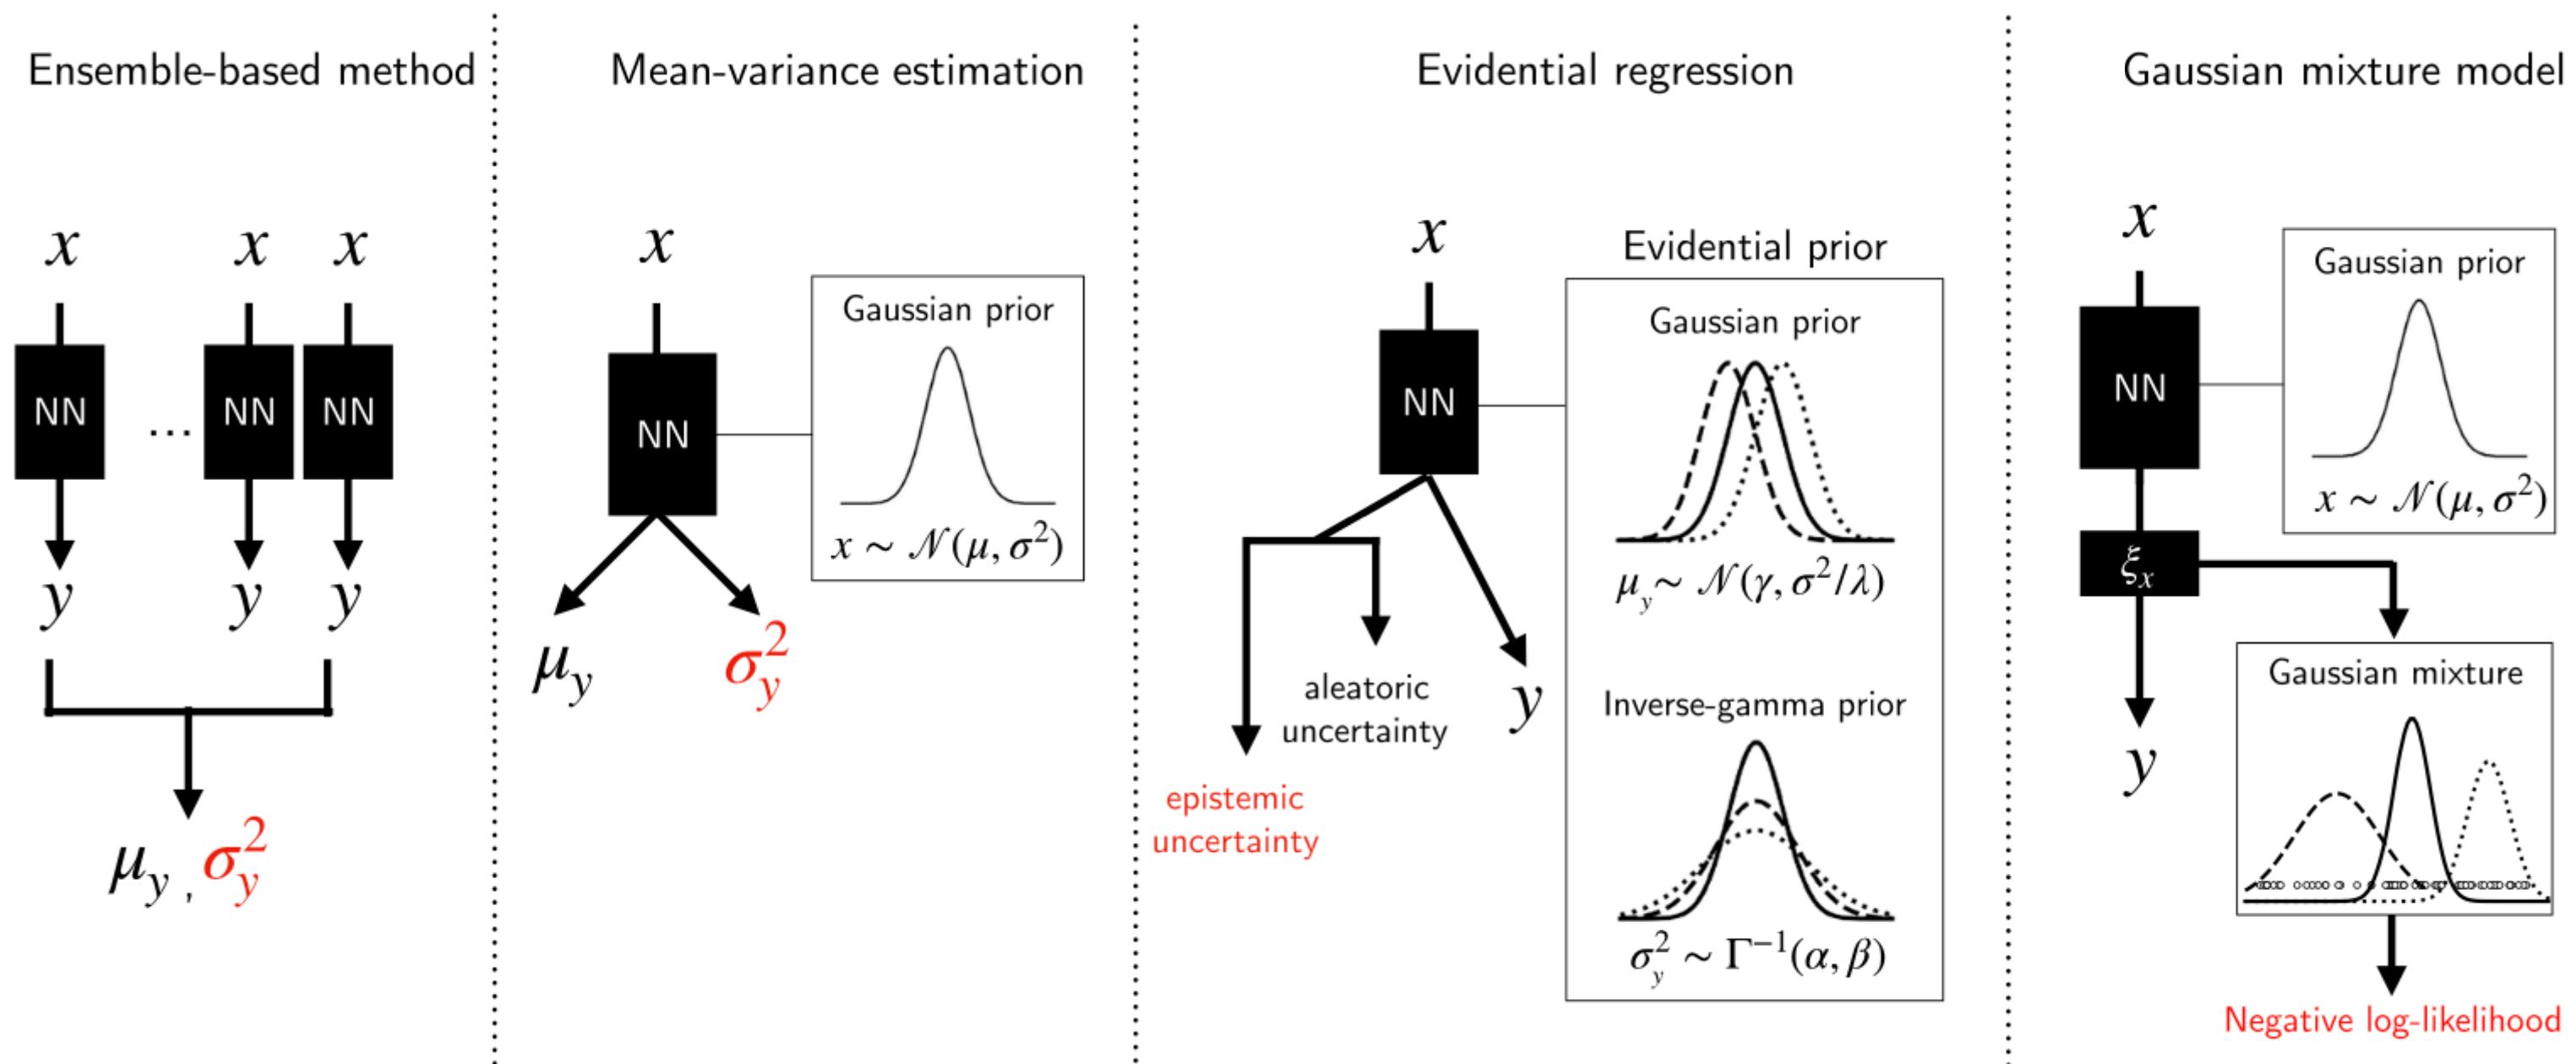
If a molecule is rotated, its energy stays the same (invariant), and its forces rotate accordingly (equivariant)

If we build stochastic networks (sampling weights) we must ensure equivariance/invariance in distribution!



Weiler, Maurice. "Equivariant and coordinate independent convolutional networks: A gauge field theory of neural networks", 2024.

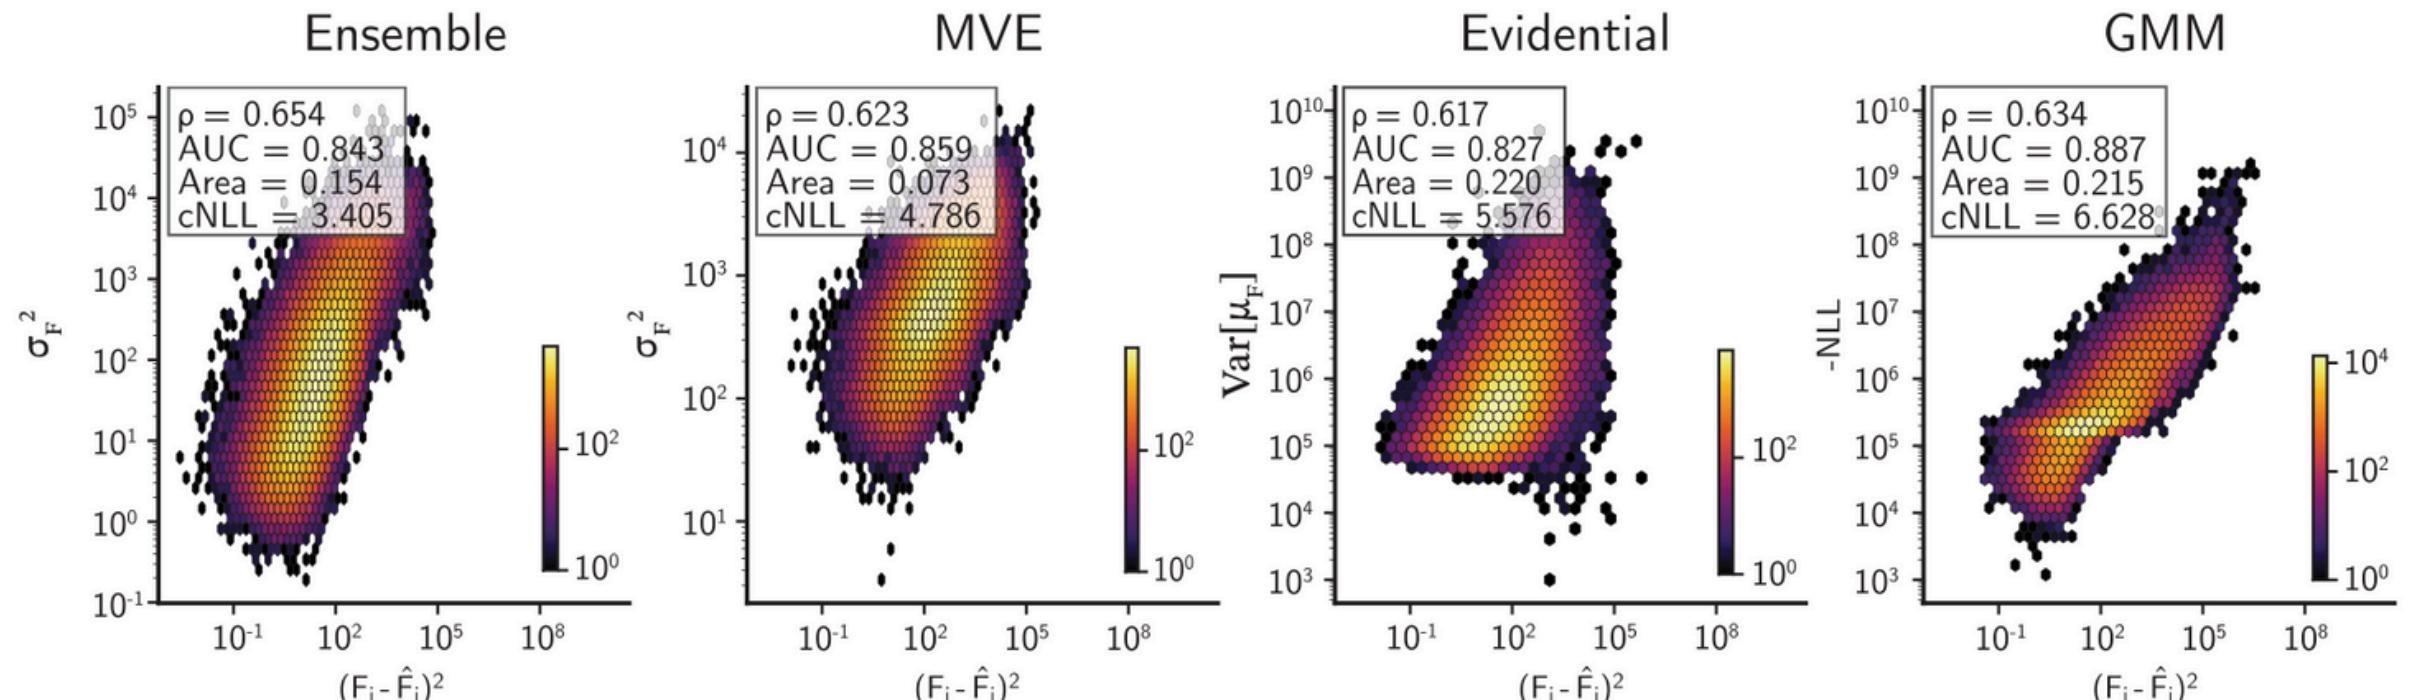
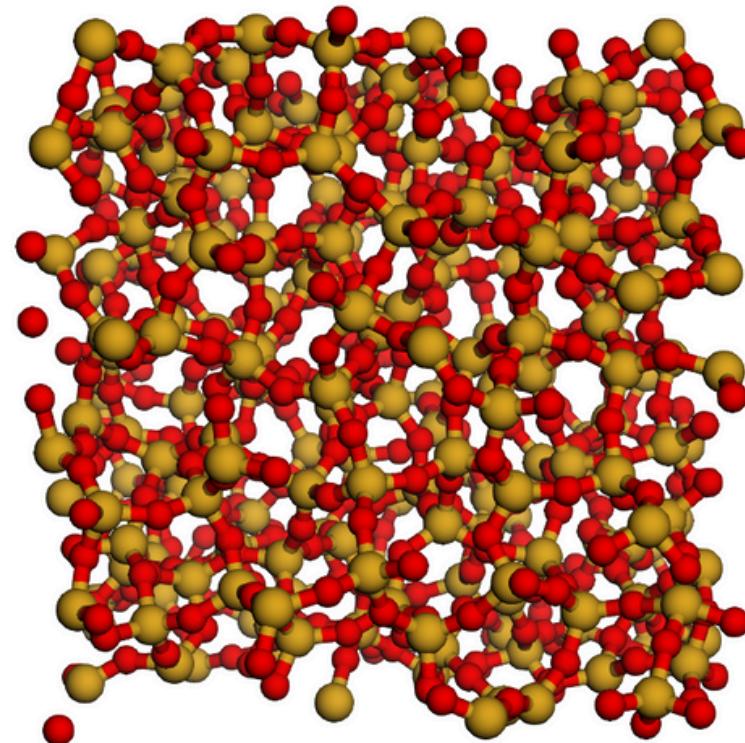
# MODELLING UNCERTAINTY IN MLFFS



Tan, Aik Rui, et al. "Single-model uncertainty quantification in neural network potentials does not consistently outperform model ensembles." *npj Computational Materials* 9.1 (2023): 225.

# MODELLING UNCERTAINTY IN MLFFS - SILICA

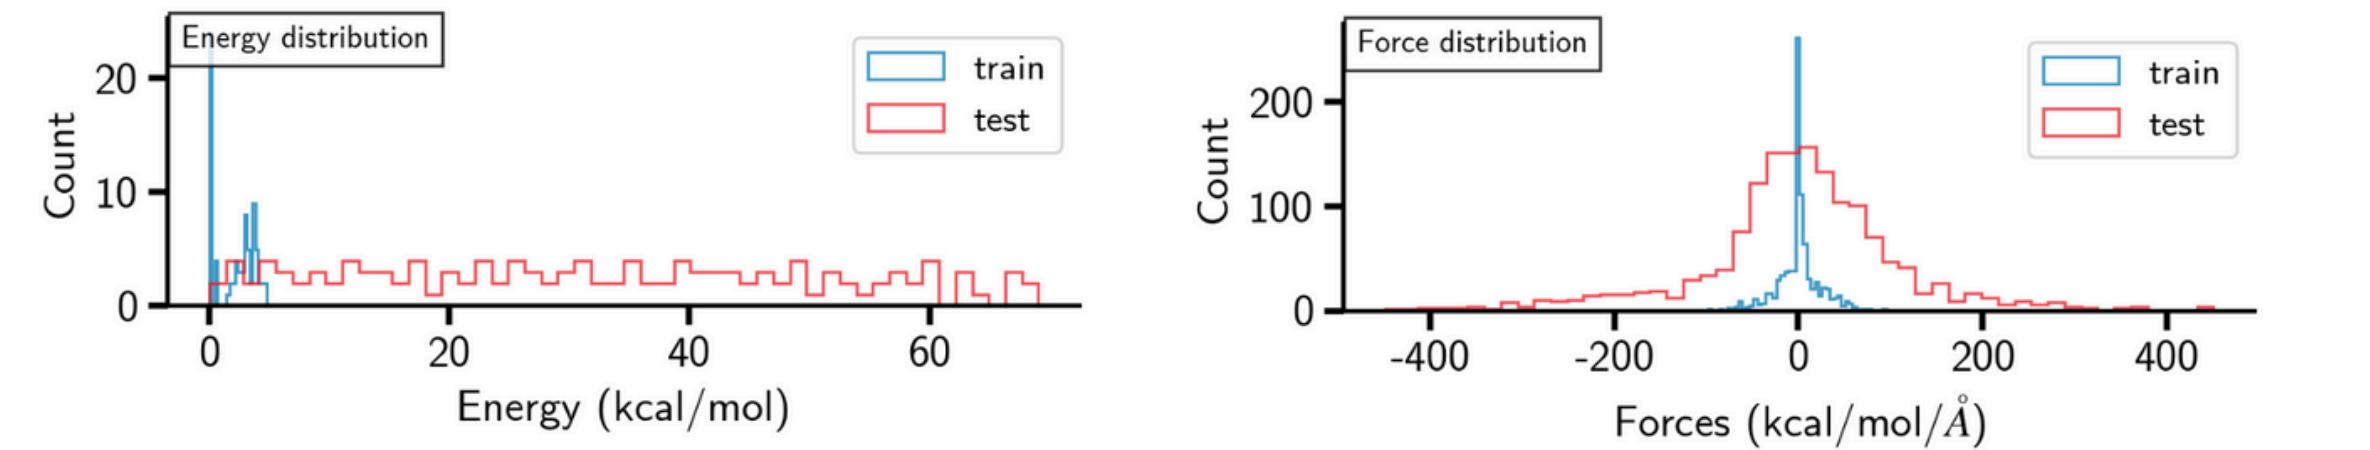
**Silica glass** is made up of large bulk structures (many atoms) and the computational cost of performing quantum chemistry calculations on these bulk systems is prohibitively high! **MLFFs can help mitigating this problem, but reliable uncertainties are needed!**



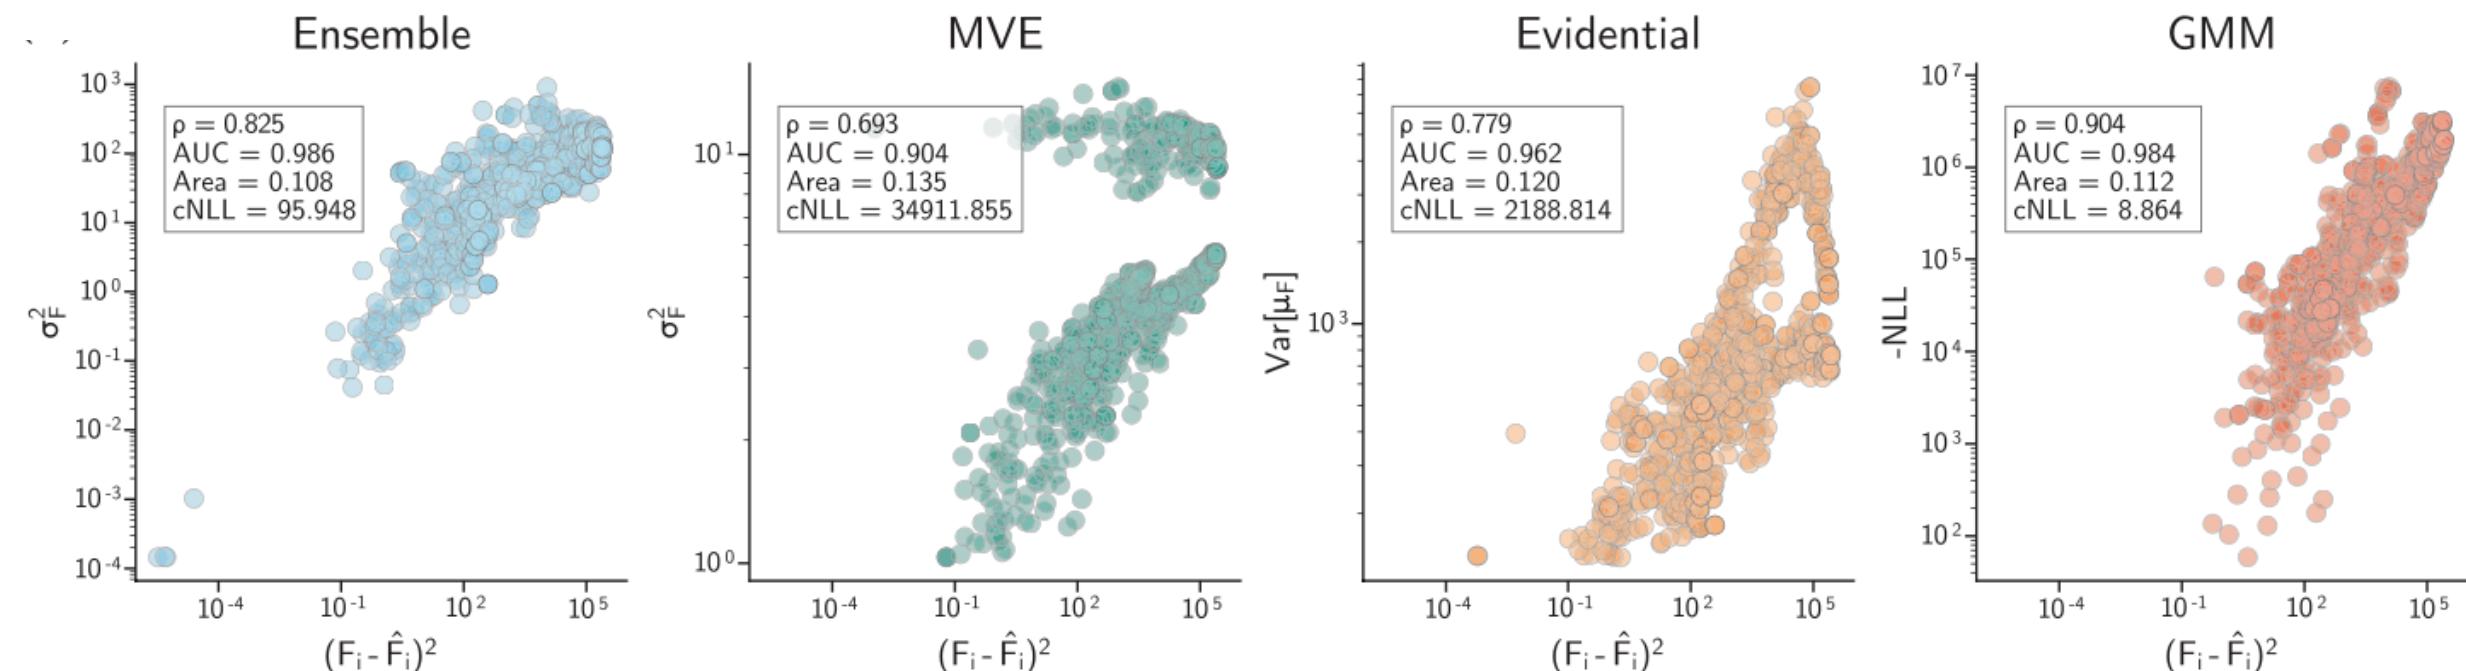
# MODELLING UNCERTAINTY IN MLFFS - AMMONIA

**Ammonia NH<sub>3</sub>** is a simple molecule with a very interesting property: ***the nitrogen atom flips from one side of the plane formed by its substituents to the other***, like an umbrella turning inside out

***Testing on heavy OOD data***

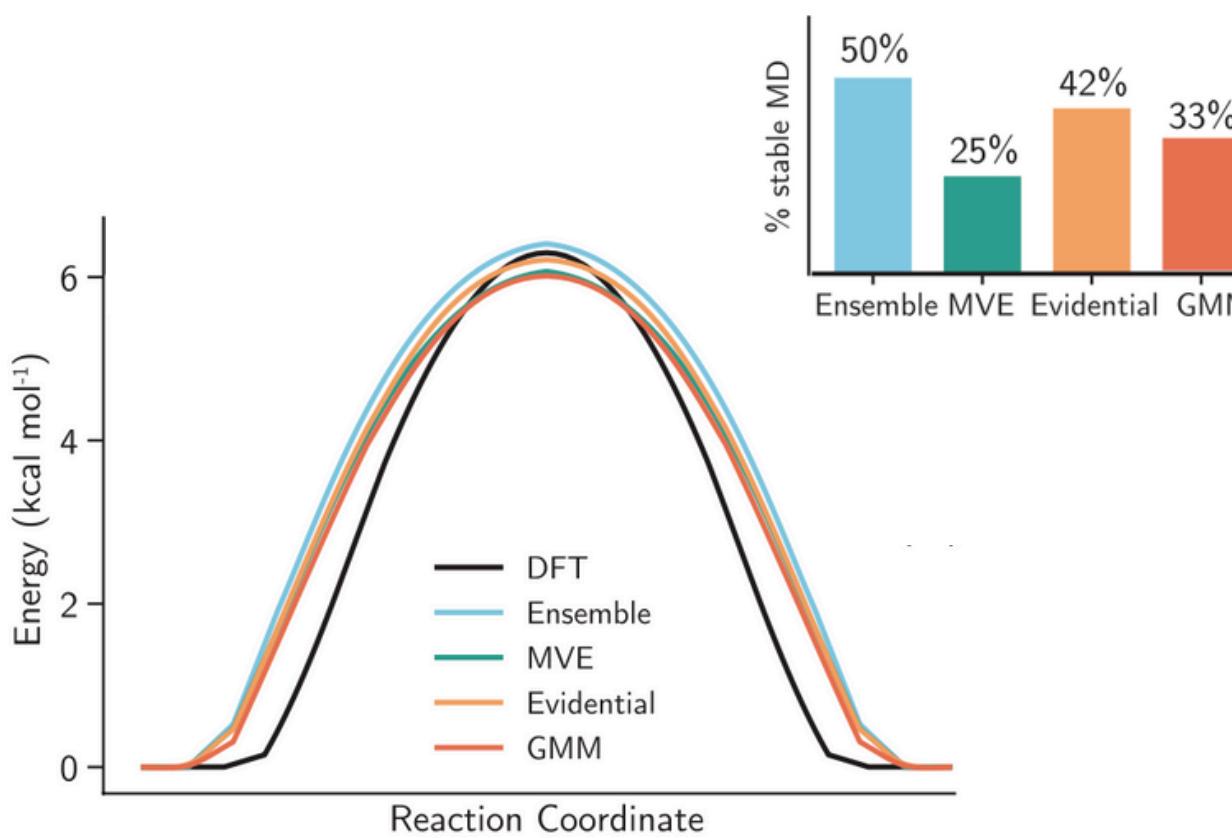


***Both Deep Ensemble and GMM perform well!***

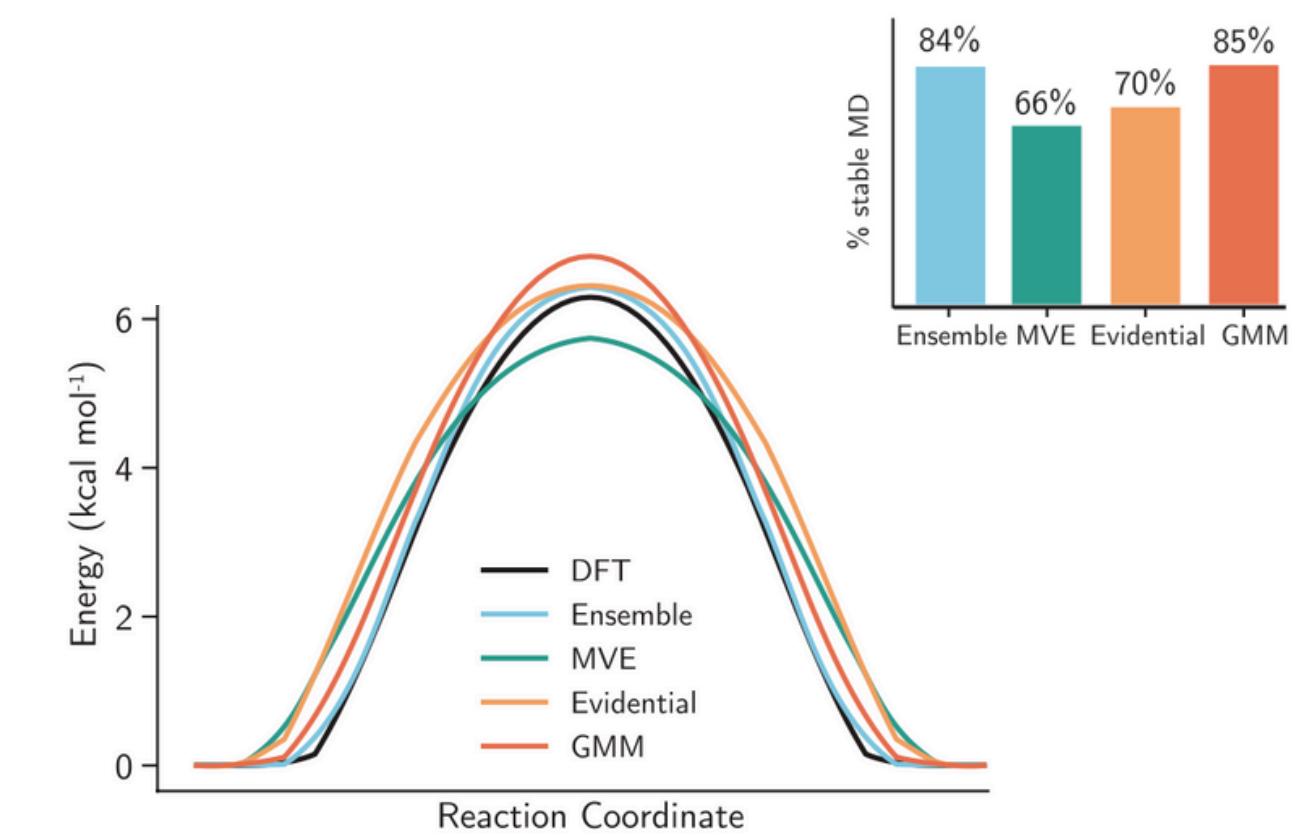


# MODELLING UNCERTAINTY IN MLFFS - AMMONIA

**Ammonia NH<sub>3</sub>** is a simple molecule with a very interesting property: ***the nitrogen atom flips from one side of the plane formed by its substituents to the other***, like an umbrella turning inside out



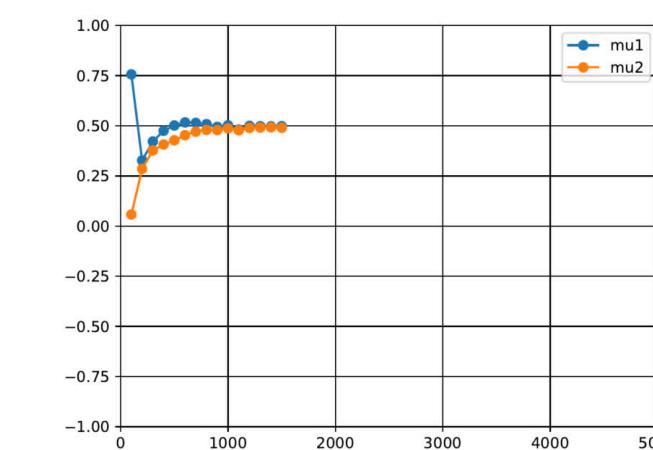
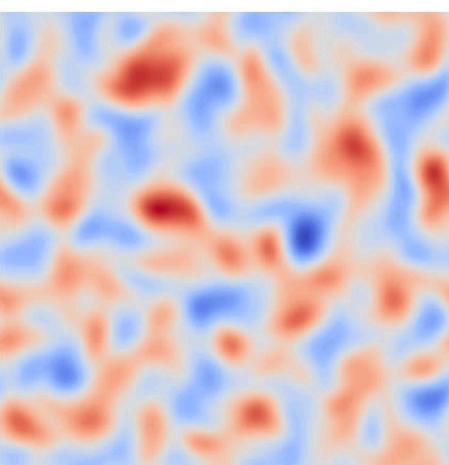
***active learning***



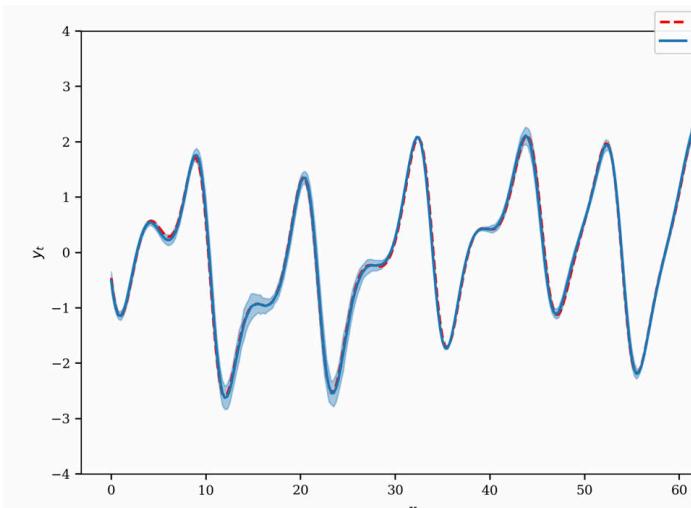
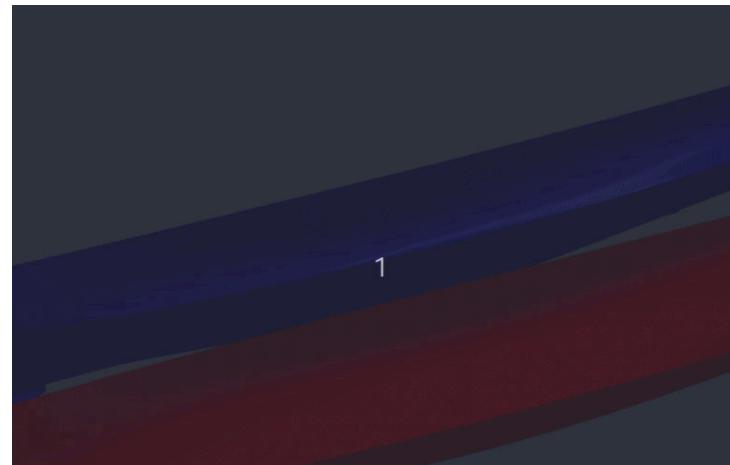
Ang, S. J., Wang, W., Schwalbe-Koda, D., Axelrod, S. & Gómez-Bombarelli, R. Active learning accelerates ab initio molecular dynamics on reactive energy surfaces. *Chem* 7, 1–32 (2021).

# CODING SOFTWARE FOR SCIML

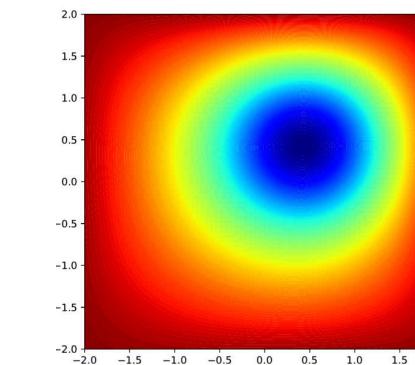
*Fast Fluid  
Dynamics  
Simulations*



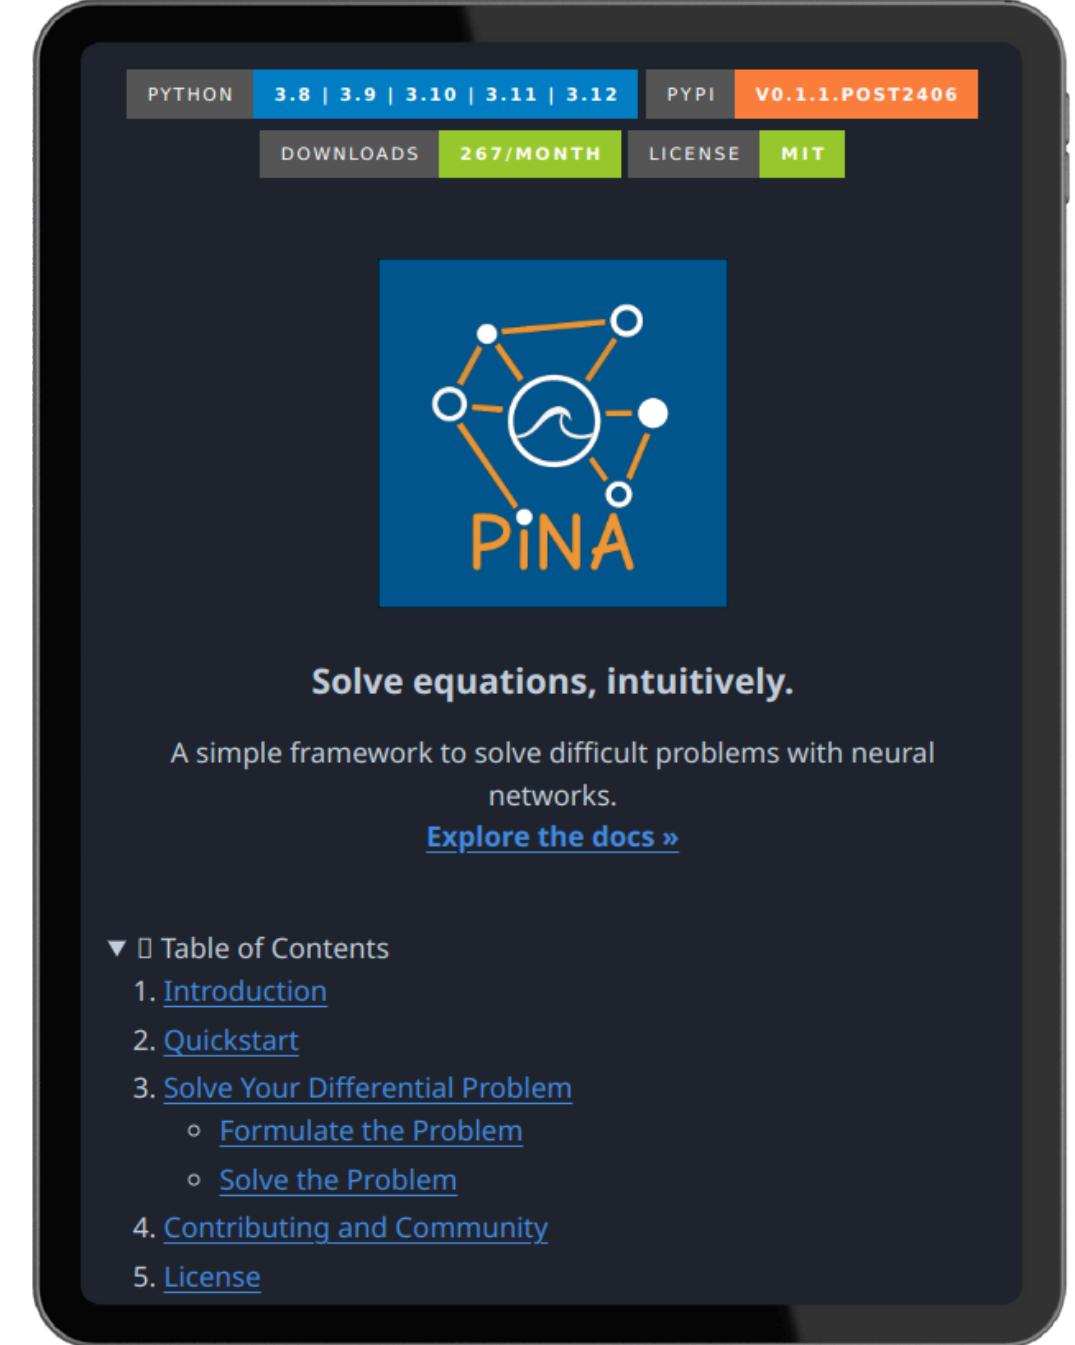
*Physics Informed  
Deformation  
Learning*



*Inverse Problems*



*Uncertainty  
Quantification  
with BARNN*



PYTHON 3.8 | 3.9 | 3.10 | 3.11 | 3.12 PYPI V0.1.1.POST2406  
DOWNLOADS 267/MONTH LICENSE MIT

**PINA**

Solve equations, intuitively.

A simple framework to solve difficult problems with neural networks.

[Explore the docs »](#)

▼ Table of Contents

- 1. [Introduction](#)
- 2. [Quickstart](#)
- 3. [Solve Your Differential Problem](#)
  - o [Formulate the Problem](#)
  - o [Solve the Problem](#)
- 4. [Contributing and Community](#)
- 5. [License](#)



Coscia, Dario, et al. "Physics-informed neural networks for advanced modeling." *Journal of Open Source Software* 8.87 (2023): 5352.